

Constructivism and Three Forms of Perspective-Dependence in Metaethics¹

KARL SCHAFER

University of Pittsburgh

1. Metaethical Constructivism

Many contemporary metaethicists are attracted to some form of normative realism. But even the most committed realists should have their doubts. For there are real epistemological and metaphysical costs to normative realism. And perhaps these costs are greater than realists have acknowledged thus far—so great that normative realism is ultimately indefensible. If this is a real possibility, then it makes good sense for realists to explore alternatives to their views.

This paper is an attempt to explore one such alternative. If it is to be attractive to someone with realist sympathies, this alternative should preserve as many of the features of normative discourse that make realism attractive as possible. Thus, it should involve a substantial element of ‘quasi-realism’. But, at the same time, it must avoid those features of realism that fuel the epistemological and metaphysical concerns about its viability.

Combining these two features is no easy task—and yet the possibility of doing so is perennially attractive.² In what follows I discuss one popular strategy for accomplishing this task: metaethical constructivism. In doing so, I will develop and defend a form of metaethical

¹ In developing this material, I have been helped by the advice of more individuals than I will be able to remember now. Particularly notable among these are Sinan Dogramaci, Jamie Dreier, Andy Egan, Max Kölbel, John MacFarlane, Mark Schroeder, Kieran Setiya, and James Shaw—as well as audiences at the Center for Human Values at Princeton University, Humboldt Universität zu Berlin, the Madison Metaethics Workshop, the University of Toronto, and my graduate seminar at the University of Pittsburgh.

² For discussion of this issue, see especially Dreier (2004).

constructivism that seems to me better suited than most to playing this role. But in order to begin exploring this form of constructivism, we need to have some understanding of what metaethical constructivism *is* and what motivates it. Unfortunately, while constructivism is one of the dominant movements within metaethics today, defining it is no easy task.³ But one thing that unites almost all constructivists is the idea that normative facts are somehow dependent upon the normative judgments or values of individuals or communities. Put in characteristic constructivist language, we may say that these judgments or values provide the ‘materials’ from which the normative facts are constructed via some ‘procedure of construction’.

In particular, contemporary constructivists have generally understood their position to imply that the normative facts about a person are the product of a constructive procedure that takes the normative judgments or values of that person (or her community) as its inputs. If we use the phrase ‘normative perspective’ to capture those features of a person’s normative judgments, values, and the like that the constructivist takes to form the relevant ‘materials of construction’, then we can capture the characteristic claim of these traditional forms of constructivism as follows:

Subject-Dependence of Normative Truth: The truth-value of a normative judgment about an individual is determined by the normative perspective of that individual (or her community).⁴

This claim is, I think, best understood as a claim about what *grounds* the truth-value of normative judgments—that is, as the claim that their truth-value is (at least in the relevant respects) grounded in facts about the normative perspective of their subjects (or those subjects’ communities). In other words, I take this to be a claim about what makes normative judgments true or false.⁵ So, for example, a constructivist of

³ For perhaps the most prominent Kantian form of metaethical constructivism see Korsgaard (1996). For a more Humean variant of this sort of view, see Street (2008). But, of course, there are many other such views in the literature. Note that certain forms of constructivism involve substantive normative constraints on the constructive procedure. As I note below, such views lie somewhere along a spectrum between full-fledged constructivism and normative realism. I discuss such views further in Schafer (forthcoming a).

⁴ Of course, it will not be determined by this perspective *alone*—further non-normative facts will also be relevant here.

⁵ For a classic statement of truthmaker theory, see Armstrong (1997). For an argument for truthmaking without *truthmakers*, see Lewis (2001) and Melia (2005). If one adopts the former line, the relevant truthmakers will be complicated combinations of the subject’s values, judgments, and desires. On the latter sort of view, normative judgments will be made true by facts about these matters. I remain neutral about these debates here.

this sort about practical reasons might believe that judgments about what a person has most reason to do are made true by facts about their values, judgments, and desires—once these are processed via the relevant ‘procedure of construction’.

Importantly, the Subject-Dependence of Normative Truth can be made precise in two quite different ways. First, the Subject-Dependence of Normative Truth might simply be a consequence of the manner in which the *content* of a normative judgment is determined by who its subject is. In this case, the truth-value of a normative judgment would be grounded in facts about the subject of that judgment because its content would be relativized to the subject’s normative perspective. This is the most common way of developing such a position, but we might also think of the relevant form of subject-sensitivity in terms of the idea that the truth-conditions of these judgments, but *not* their contents, are grounded in facts about their subject. In this case, the dependence of a normative judgment’s truth-value on its subject’s perspective would not be due to its content, but would instead be a matter of the manner in which the truth-value of this content is determined.

For the moment, little will depend on which of these versions of the Subject-Dependence of Normative Truth we accept. Rather, what interests me here is simply this principle itself—which is common currency among most contemporary constructivists.

What constructivists *do* disagree about is the degree to which the relevant procedure of construction constrains the normative facts that result from it. In other words, contemporary constructivists take very different views about whether the fact that the truth-conditions of a normative judgment are *grounded* in the subject’s normative perspective means that these truth-conditions will be *sensitive* in interesting ways to the subject’s *particular* normative perspective.

So-called ‘Humean constructivists’ believe that there are no non-trivial constraints on the values that can be incorporated into a coherent normative perspective.⁶ Moreover, they take the relevant procedure of construction to involve nothing beyond the purification of an individual’s state of mind so as to achieve this sort of minimal coherence. As a result, they believe that the relevant procedure of construction places no substantive constraints on the normative facts that are constructed through it. Start with the right (or wrong) materials of construction, they claim, and the relevant procedure of construction will generate nearly any result you wish. Thus, on this view, if we accept the Subject-Dependence of Normative Truth, we must also

⁶ I take this terminology from the discussion of these issues in Street (2008).

accept that the truth-value of a normative judgment is highly sensitive to facts about its subject's normative perspective.

The consequences of this view are often made explicit via consideration of individuals such as an 'ideally coherent Caligula', who highly values inflicting suffering on the innocent and nonetheless is completely coherent in his values and beliefs.⁷ If such individuals are possible, and the procedure of construction is the simple one just described, then one consequence of constructivism about practical reasons will be that such persons have every reason to behave in a horribly sadistic fashion. On such a view, so long as a sadist is fully coherent in his sadism, we have no basis for denying that he has good reason to torment and torture to his heart's delight.

Not surprisingly, such consequences have led many philosophers to reject *pure* Humean versions of constructivism. One possible response to these concerns is to introduce *substantive* normative constraints into the constructivist's characterization of the materials or procedure of construction.⁸ So, for example, we might build some sort of universality requirement into our conception of the relevant procedure of construction so as to rule out the results that seemed so troubling above. This maneuver has much to recommend it.⁹ But if these additional elements are taken by the constructivist to be *brute* constraints on how the normative facts are constructed, they introduce a significant element of normative realism into the constructivist's understanding of the procedure of construction. And this alone will make views of this sort unattractive to many constructivists, who are motivated precisely by a desire to avoid an appeal to brute normative facts or constraints of this sort.

For this reason, many contemporary constructivists have attempted to ground further constraints of this sort in the very nature of agency or practical thought itself. According to such "Kantian constructivists", some meaningful constraints on what an agent can coherently value are built into the very nature of agency or practical reasoning—constraints that make it the case that even Caligula has good reason to treat other human beings with respect. Thus, such views deny the Humean's inference from the claim that the truth-value of normative judgments is grounded in facts about the subject to the claim that the truth-value of such judgments can *vary* radically with differences in the subject's normative perspective.

⁷ Or a community of such individuals in the case of non-individualistic forms of constructivism.

⁸ A view of this general sort is perhaps not implausibly attributed to Kant. For an important contemporary view in this neighborhood, see Schroeder (2007).

⁹ See again Schafer (forthcoming a) for further discussion.

This is an extremely important style of metaethical constructivism, and it has quite rightly been the focus of much of the recent debate concerning constructivism. But, without rehearsing the details of these debates, there is a good deal of skepticism in the literature about whether the Kantian constructivist can demonstrate that sufficiently powerful constraints of this sort are built into the very nature of agency or practical thought in the manner this view requires.¹⁰ So while I do not think the prospects for this form of constructivism are entirely bleak, it is hard to feel confident that this avenue for defending constructivism will ultimately defang every objection in this general class.

If this is right, it can easily seem as if traditional forms of constructivism face a choice between two alternatives that are equally unattractive, even by the constructivists' own lights. Either they may remain pure in their constructivism and accept that their view does have certain very counterintuitive consequences in cases like the ideally coherent Caligula,¹¹ or they may introduce substantive normative constraints into their characterization of the 'materials' and 'procedure' of construction—constraints which are not the product of the nature of action or agency—in virtue of which they will take on many of the burdens associated with normative realism.

2. From Constructivism to Perspectivalism

Of course, none of these objections on its own shows that traditional forms of constructivism should be rejected. But they do give those who are sympathetic to these views cause to explore other ways of developing the basic ideas behind metaethical constructivism.

In order to consider the different ways these ideas might be developed, it is helpful to begin by considering what motivates metaethicists to accept a form of constructivism in the first place. Naturally enough, there are a number of different considerations that one might cite in response to this question; but perhaps the central motivation for constructivism is given particularly clear expression in the following passage by Sharon Street:

This I regard as the fundamental metaethical insight of ... constructivist views—the idea that determinations of the correctness of normative judgments must always be made *from somewhere*—and in particular, from some practical point of view, constituted by the acceptance of

¹⁰ For different variants of these worries see Enoch (2006), Setiya (2007), and Street (2008).

¹¹ For such a view, see Street (2008).

further normative judgments. From nowhere, according to constructivist views in ethics, there are no normative facts.¹²

Street—like many contemporary constructivists—takes this thought to support the Subject-Dependence of Normative Truth. But the basic idea here is considerably more general than this—namely, that normative truth must be understood as perspective-dependent in *some* way.¹³ That is, Street’s comments directly support at most the following claim, put (as above) in terms of grounding:

Perspective-Dependence of Normative Truth: The truth-value of a normative judgment about an individual is grounded (at least in the relevant respects) in facts about *some* relevant normative perspective.¹⁴

But this general idea leaves unspecified exactly *which* perspective should be taken to be relevant to the truth-value of a particular normative judgment. Here there are at least three salient possibilities:

1. **Subject-Dependence of Normative Truth:** The truth-value of a normative judgment about an individual is grounded (at least in the relevant respects) in facts about the normative perspective of that individual (or her community).

¹² Street (2008, p. 14).

¹³ For Korsgaard, at least, the Subject-Dependence of Normative Truth is motivated by a further consideration: namely, the idea that facts about what counts as a reason for some agent should be capable of engaging *that individual* from within his or her own normative perspective. Unfortunately, this constraint itself *seems* to clash with some of our ordinary, considered normative judgments. Or at least it does unless we can derive robust constraints on what an individual’s reasons for action are from the nature of agency itself in the manner Korsgaard claims we can. After all, when we consider an individual like the ideally coherent Caligula, we feel entirely comfortable with the claim that Caligula is acting in a manner that is contrary to what he ought to do, even though we believe that this verdict will not resonate in any way with his own normative perspective. Thus, if our goal is to preserve as much as is possible of the surface of normative discourse, we should be suspicious of any argument for the Subject-Dependence of Normative Truth that *begins* with the claim that facts about an individual’s reasons (etc.) *must* always be rooted in his or her own normative perspective. To put things in Korsgaard’s terms, while we should expect to be able to give an account of the normative facts that answers the ‘normative question’ from our own point of view, we should not expect to be able to give an account of these facts that will answer this question from every possible point of view—even those that are profoundly unreasonable. So, even if this philosophical ambition is in fact defensible along the lines Korsgaard suggests, I am skeptical that it is part of our *ordinary* understanding of normative discourse.

¹⁴ For reasons that will become clear in a moment, the reference to ‘truth’ here is importantly ambiguous.

2. **Judger-Dependence of Normative Truth:** The truth-value of a normative judgment about an individual is grounded (at least in the relevant respects) in facts about the normative perspective of the individual making the judgment (or her community).
3. **Assessor-Dependence of Normative Truth:** The truth-value of a normative judgment about an individual is grounded (at least in the relevant respects) in facts about the normative perspective of the individual assessing the judgment (or her community).¹⁵

Once again, each of these forms of perspective-dependence can be understood either in terms of the idea that the contents of normative judgments are perspective-dependent or in terms of the idea that the truth-conditions, but *not* the contents, of such judgments are so dependent. In the case of the Judger-Dependence of Normative Truth, the first option will generate a form of what is sometimes called “indexical contextualism”—while the second will generate a form of so-called “non-indexical contextualism”.¹⁶ Once again, since my main aim here will be to argue for *some* form of the third of these options, I won’t focus on the question of which of these two ways of making these views precise is preferable here—but, for reasons I discuss elsewhere, it is not always easy to see what substantial issue separates sophisticated forms of non-indexical contextualism from views that accept the Assessor-Dependence of Normative Truth.¹⁷ So, to some degree, I want to allow that certain sophisticated forms of non-indexical contextualism could share most of the costs and benefits of the view discussed below.

What I do want to focus on is the ability of these different views to make sense of our ordinary pattern of considered normative judgment. In this respect, the crucial difference between these three views is that the first two appear to require someone who is assessing the truth-value of a normative judgment to defer to another normative perspective in making that assessment—be that the perspective of the subject of that judgment or of the person who made it.¹⁸ The Assessor-Dependence of

¹⁵ The phenomenon of assessor-dependent truth-conditions is the focus of much recent work in the philosophy of language. For a prominent example of such work, see MacFarlane (2005). Kölbel (2002) and Egan (2007) are also important here, although their views are more difficult to characterize.

¹⁶ For this distinction, see MacFarlane (2009).

¹⁷ See Schafer (forthcoming c).

¹⁸ Again, one might contest this point with respect to certain sophisticated forms of non-indexical contextualism. My own view is that the forms of non-indexical contextualism that escape these worries end up being mere notational variants of the position I will be arguing for here.

Normative Truth, on the other hand, does not require this, insofar as it allows a single normative judgment's truth-value to vary with the individual who is assessing it.

As just noted, constructivists have tended to understand the claim that normative truth is Perspective-Dependent in terms of the idea that it is Subject-Dependent. But, as we just observed, this raises difficult issues about whether such a view is able to preserve our considered normative judgments when it comes to individuals like an ideally coherent Caligula. After all, if we think of the truth of a normative judgment as Subject-Dependent, then when I make normative judgments about other persons, I will have to do so in a manner that is responsive, not to my own normative views, but to theirs. And this, at least *prima facie*, will force us to be more deferential in our normative judgments than we often find intuitive.

Thus, if we hope to develop a form of constructivism that is maximally faithful to these intuitions, we have good reason to consider the other two forms of Perspective-Dependence just noted. Unfortunately, the second such possibility—a form of metaethical *contextualism*—suffers from similar problems.¹⁹ For, on this view, when we consider whether someone else's normative judgments or assertions are true or false, these assessments should be responsive not to our own normative views, but instead to the normative views of the person who made the judgment under consideration. So, for example, when the ideally coherent Caligula *judges* that he has good reason to act sadistically, we will find ourselves forced, on this view, to take this judgment to be true. Once again, this conflicts with our considered judgments about these cases—and, in particular, with the sort of objectivity that normative talk and thought intuitively appears to possess.²⁰ Thus, neither of these views appears to do full justice to our considered judgments—be they our first-order normative judgments, as in the first case, or our judgments about the truth and falsity of such judgments, as in the second.

Fortunately, we can do better in these two respects by thinking about perspective-dependence in the third way noted above—so that the truth-value of a judgment is grounded (in part) in facts about the perspective of the person assessing its truth. In order to take this path, given that

¹⁹ Views of this sort are often referred to as forms of “speaker relativism”. For an excellent example of this sort of view, see Dreier (1990).

²⁰ This worry is simplest to generate about standard forms of contextualism—where the perspective of the judger determines the *content* of her judgment. The issues here are more complex with respect to “non-indexical” forms of contextualism—especially if we take the assessment of a judgment's truth-value *always* to involve a further judgment about that judgment. For an excellent defense of this form of contextualism in the moral case, see Brogaard (2008).

this implies that the truth-conditions for such judgments will be sensitive to differences between assessors, we will want to adopt a formal semantics for normative judgments and assertions according to which the truth-value of a judgment or assertion can vary with the perspective from which it is assessed.²¹ As I will discuss shortly, there are a number of serious worries about any view of this sort. But *if* these worries can be answered, a view of this sort would do far better than the two just discussed when it comes to preserving our ordinary judgments in cases like the ideally coherent Caligula. For, given this last form of perspective-dependence, when we assess whether someone else's judgments are true or false, this assessment ought to be responsive to *our own* normative views—and not the views of the subject of the judgment or the person who made it. In this way, thinking of Perspective-Dependence in terms of Assessor-Dependence holds out the prospect of preserving much more of the superficial objectivity of normative discourse than either of the other two possibilities just discussed.

In considering this point, it is particularly important to remember that when someone who endorses the Assessor-Dependence of Normative Truth goes on (as I've suggested he should) to claim that the truth of normative judgments is sensitive to the perspective of the assessor, he is not making a claim about the truth-value of normative judgments in the ordinary sense of these terms. Rather, he must be understood to be making a claim about their truth-value in the technical sense of those terms with which one works when doing formal semantics.²² For our ordinary concept of truth obeys some variant of the equivalence scheme that allows one to move freely between the judgment that P and the judgment that P is true. And this scheme does not even make sense once we think of a judgment's truth as sensitive to the perspective of the assessor. Thus, someone who thinks of the truth of normative judgments as sensitive to the perspective of the assessor must also think of our *ordinary* judgments of truth and falsity as sensitive to the perspective of the assessor when those judgments concern normative judgments. For we can only preserve the equivalence between the judgment that P and the judgment that P is true (in the ordinary sense), where

²¹ In what follows, I will discuss the semantics of normative judgments and assertions in parallel with one another. This is possible only insofar as both thought and public language are the sorts of things it makes sense to provide a formal semantics for. But if one is uncomfortable about this assumption, one can easily take my comments here to be limited to the case of public language.

²² On some views these two concepts of truth will be one and the same. But any such identification greatly limits the resources available to the semanticist—especially if we conceive of the ordinary notion of truth in minimalist terms. As a result, my sense is that few contemporary philosophers or linguists would accept it without qualification.

the truth of P is assessment-sensitive, by taking the judgment that P is true (in the ordinary sense) to be assessment-sensitive in the same way as the judgment that P .

Thus, the assessor-dependent theorist must distinguish between a technical notion of truth— truth_S —and the ordinary notion of truth— truth_O —that we all employ. When this theorist provides his formal semantic theory, he will do so in terms of truth_S . And it will be with respect to this notion that the truth_S of a normative judgment can only be evaluated in an assessor-sensitive fashion. On the other hand, this formal semantic theory will treat claims about the truth_O of a judgment so that the truth conditions of these claims are the same as the truth conditions of the judgment they are concerned with. Thus, the truth_S of the claim that some normative judgment is true_O will be assessment-dependent in just the same way as the truth_S of the underlying normative judgment is.

Given all this, someone who accepts the Assessor-Dependence of Normative Truth should take this claim, in the first instance, to involve a claim about what grounds the truth_S -conditions for normative judgments—namely, that these facts are grounded in facts about the normative perspectives of the individuals who are assessing those judgments' truth-value. But given this, what is the status of truth_O -conditions on this view? One option here would be to claim that the truth_O -conditions of a normative judgment are in turn grounded in that judgment's truth_S -conditions. After all, on the view being developed here, whether or not some normative judgment is true_O will itself vary with the normative perspective from which this question is assessed. Thus, it might be natural for someone who accepts the Assessor-Dependence of Normative Truth to think of the truth_O and falsity $_O$ of normative judgments as a straightforward product of their assessment-sensitive truth_S -conditions relative to the relevant contexts of assessment. For example, one might want to claim that the fact that some normative judgment J is true_O is explained by the fact that J is true_S relative to the relevant context of assessment.

I think a view of this general sort has something to recommend it, but I know from conversation that many philosophers find the explanatory claims made here mysterious. Thus, it might be better for our theorist simply to insist that, on his view, there is no further substantive fact about truth_O -conditions to explain.²³ For instance, he might insist that on his view all there really is to the structure of normative discourse is the truth_S -conditions of normative judgments and assertions.²⁴

²³ Compare Sider's discussion of non-substantiality in his (2011).

²⁴ For more discussion of these issues, see Schafer (forthcoming b).

I won't choose between these views here. Rather, what is most important here is the fact that, on an assessment-dependent view, so long as we restrict ourselves to our ordinary notions of truth, fact, and the like, our assessments of other person's normative judgments will always be made from within our normative perspective. As a result, these assessments will fail to make explicit the perspective-sensitivity that, according to the Assessor-Dependence of Normative Truth, is present in these cases. Rather, as long as we restrict ourselves to these sorts of assessments, normative discourse will appear to be fully objective. It is only when we turn to the more technical vocabulary of the formal semanticist (i.e. to truths_S as opposed to truths_O), and to related claims about grounding, that the perspective-dependence of normative truth becomes clear. In other words, unlike the other forms of perspective-dependence noted above, the perspective-dependence involved in an assessment-sensitive account of normative semantics leaves the non-relativist surface of *ordinary* normative discourse intact.

Of course, this means that the theorist who accepts the Assessor-Dependence of Normative Truth can only distinguish his view from normative realism in terms of the semi-technical notion of truths_S. And this raises concerns, familiar from recent discussions of quasi-realism, about whether the resulting view is anything but a notational variant of realism. In a moment, I'll return to this issue. But, by way of foreshadowing, we may note that a plausible way of responding to this worry is to show that an assessor-sensitive formal semantics has normative consequences that diverge from those of realism.

In any case, while the Perspectivalist's view only becomes clear when we turn our attention to the more technical notions of truths_S and falsity_S in which it is stated, we should not forget that it is nonetheless a formalization of a quite intuitive picture of the nature of normative judgment. As the appeal of constructivism makes clear, the idea that normative assessments are perspective-dependent in some sense is not unattractive. And yet, as we have already discussed, the ways of making this idea precise given by the Subject-Dependence or the Judger-Dependence of Normative Truth appear to conflict with many of our considered normative judgments. The explanation for this, one might suggest, is that while we do ordinarily take normative judgments to be perspective-sensitive in some manner, it is neither the perspective of the subject nor that of the judger that fixes the truth-value of a normative judgment. Rather, one might argue, our intuitive understanding of the perspective-dependence of normative truth involves the idea that a normative judgment may, in principle, be assessed from any coherent normative perspective—yielding different verdicts about its truth-value depending on the perspective from which it is assessed.

It is just this way of thinking about perspective-dependence that I am interested in formalizing here. Of course, in doing so, we must use certain technical notions. But we should not pretend that this means that there is no connection between these notions and our ordinary way of thinking about these matters. For, in fact, there is a case to be made that this sort of view involves the most straightforward way of formalizing the mixture of ‘superficial realism’ and ‘deep perspective-dependence’ that is involved in many ordinary people’s understanding of normative truth.

Putting these issues aside for a moment, what we can say at this stage is that the prospects for a form of constructivism that preserves our considered judgments about individuals like the ideally coherent Caligula are strongest if we conceive of the Perspective-Dependence of Normative Truth in assessor-dependent terms. This should, I think, come as no great surprise, given the manner in which theorists like Street have characterized the ambitions behind constructivism. After all, Street’s claim that “the determination of the correctness of a normative judgment” must be made “from somewhere” leads quite naturally in this direction. Street’s idea is that when assessing the truth of normative judgments, we must always do so *from some normative perspective*. But if this is so, why should someone making such assessments always defer to a normative perspective other than his own when doing so? Of course, sometimes we may want to evaluate a normative judgment from someone else’s perspective, but surely the most natural perspective from which to assess a normative judgment is just the normative perspective one actually occupies in assessing it?

Moreover, if we think of the constructivist’s project in this way, he can avoid making the difficult choices that the arguments of the last section attempted to force upon him. For so long as we evaluate Caligula from within our own normative perspective, the Assessor-Dependence of Normative Truth will not generate counterintuitive results. In particular, this view will not imply that Caligula ought to act in a sadistic manner—nor will it imply that Caligula’s claims about how he ought to act are true in the ordinary sense of the term (true_O). At most, the Assessor-Dependence of Normative Truth will imply that *when Caligula’s views are assessed from his own normative perspective* it is true_S that he ought to do the hideous things he believes to be right and good. And this (semi-technical) result is hardly as counterintuitive as the results above—if indeed we have any intuitions about it at all.

It would not be unnatural to refer to the resulting view as ‘metaethical relativism’—but, since this label is already taken several times over, I will refer to it instead as metaethical perspectivalism, since this

captures the radical sort of perspective-dependence involved viewing the normative domain in assessor-dependent terms.²⁵

Metaethical Perspectivalism: A normative judgment can only be assessed as true or false from within a normative perspective. In particular, the truth_S of a normative judgment about an individual is grounded (in the relevant respects) in facts about the normative perspective of the individual assessing the judgment (or her community).

To flesh out the perspectivalist's view, let's briefly consider its application to a particular case. Say we have some individual (or community)—call him Fred—and we want to know whether some normative judgment or assertion is true when assessed from Fred's normative perspective. In order to identify this perspective, consider Fred's current state of mind. Then idealize this state of mind via a procedure of reflective equilibrium until it is internally coherent and consistent. (The resulting mental state will normally be somewhat indeterminate—since there will be more than one result of this process.) Finally, extend the resulting state of mind so that it includes all of the logical, conceptual, and instrumental implications of the mental states contained within it. This will generate a (somewhat indeterminate) state of mind which represents what Fred's state of mind would be under conditions of full logical, instrumental, and conceptual rationality.²⁶

Now, for the particular normative judgment in question, examine the relationship of this judgment to the resulting state of mind. If the judgment is (determinately) incompatible with this state of mind, we may say that it is false as assessed from Fred's perspective.²⁷ If, on the other hand, the judgment is (determinately) part of this state of mind, then we may say that it is true as assessed from Fred's perspective. And if

²⁵ Brogaard (2008) also uses the label “perspectivalism” for her form of “non-indexical contextualism”—which, as noted above, is fairly closely related to the view I am defending here. Although I return to the relationship between these two views below, there isn't space to discuss this issue in detail. In any case, I hope my use of this label does not seem disrespectful.

²⁶ Thus, a minimal conception of logical, instrumental, and conceptual rationality will be built into the way the perspectivalist conceives of normative perspectives. Whether or not this is problematic for the perspectivalist is a complicated issue, but I take it that status of these minimal aspects of rationality are at least less problematic than the status of other more substantive normative demands. And, in any case, it appears that any story about how conceptual or inferential role determines truth-conditions will have to appeal to facts of this general sort.

²⁷ The precise nature of the compatibility at issue here will depend somewhat on one's views concerning the nature of normative judgments. I leave these issues to the side here, although they are of great import.

neither of these is determinately the case, we may say that it is indeterminate in truth-value relative to this perspective. This story, of course, leaves many of the details of the perspectivist's view to be determined. But ultimately this strikes me as an advantage—for, as we will begin to see in a moment, this sort of view may be developed against the background of quite different positions in the philosophy of language.

In any case, in considering the perspectivist's view, it will surely not have escaped some readers that, in moving from Humean and Kantian forms of constructivism to perspectivalism, we are traveling along a path that is familiar from the history of philosophy. For it is at least arguable that certain prominent views in post-Kantian philosophy take the form of very much this response to Kant. For example, it is not unnatural to view Nietzsche as endorsing metaethical views that represent a perspectival reaction to Kant of this sort. Thus, just as it is common to speak of 'Humean' and 'Kantian' forms of constructivism, we might use the label 'Nietzschean constructivism' to refer to the perspectivist's view.

To do so, though, would involve us in difficult questions of Nietzsche interpretation that are best avoided here. So I will leave this label to the side.²⁸ But this is not the only familiar metaethical view that perspectivalism should call to mind. For, as many readers will have surely noted, the resulting form of constructivism is very similar on a structural level to the sophisticated forms of expressivism that are associated with Simon Blackburn and Allan Gibbard, among others.²⁹ Thus, it is hardly surprising that—as we will see—many of the concerns that are familiar from the debate about their views also arise for the perspectivist.

That having been said, the perspectivist *appears* to differ from the expressivist in remaining committed to a thoroughly truth-theoretic (and arguably truthmaker-theoretic) conception of normative semantics. Of course, it is impossible to say how deep this difference is without a careful comparison of the perspectivist's position with specific

²⁸ For instance, Katsafanas (forthcoming) is currently developing a broadly constructivist reading of Nietzsche, which—contrary to what I have just suggested—interprets Nietzsche as endorsing a version of the Subject-Dependence of Normative Truth.

²⁹ See, for example, Blackburn (1998) and Gibbard (2003). In particular, as Dreier (2009) has observed, there is a natural isomorphism between the sorts of "fact-plan-worlds" that are the focus of Gibbard's form of expressivism and the perspectivist's assessor-dependent account of the truth-conditions of normative judgments. Thus, to some degree, the perspectivist's account of normative semantics could be viewed as a way of re-characterizing Gibbard's account of normative semantics so as to stress the continuities between his account and more traditional truth-conditional ones.

forms of expressivism. And there is no space here to take up this topic in earnest. So for the moment I want to remain content with the observation that, while perspectivalism is certainly a close cousin of views like Gibbard's, it also appears to make notions of truth and falsity more central to normative semantics than such views do.³⁰

3. Objections to Perspectivalism

If this is right, there are compelling *prima facie* reasons to think that perspectivalism is an attractive way of developing the basic constructivist idea that the normative truth is perspective-dependent. But if it is to be more than *prima facie* compelling, it must be defensible in the face of objections. And there are a number of worries about such a view that many will take to be fatal for perspectivalism:

1. As noted above, one might worry that perspectivalism will turn out to be a mere notational variant of realism, unable to support the weight of its anti-realist pretensions.
2. In the opposite direction, one might worry that perspectivalism will be unable to account for the relatively robust forms of disagreement we find in the normative domain.³¹
3. Connected to both these worries, one might wonder what it could even mean to say that the truth-value of a judgment varies from perspective to perspective. Following John MacFarlane (forthcoming a), we can call this the worry that we cannot make sense of the notion of perspective-relative truth.

In the following, I will consider the first two of these objections in turn. In doing so, I hope to sketch one way the perspectivist might respond to them. Along the way, I also hope to show where the

³⁰ I say a bit more about these issues at the close of this essay. For a particularly clear statement of Gibbard's views about the place of the concepts of truth and falsity in semantics, very generally construed, see the second Appendix to his (forthcoming).

³¹ For some of the complexities involved in this issue see MacFarlane (2007) and Dreier (2009). As both MacFarlane and Dreier emphasize, the implications of a given formal semantic apparatus for the question of when and how individuals disagree are far from obvious. Thus, there is plainly a real need—which should, I think, be felt equally by all parties to these disputes—for a systematic account of the different forms of disagreement and their relationship to one another. (Dreier also quite correctly stresses the connections between these issues and concerns about the status of negation in these systems. For the sake of simplicity, I focus on what is likely to seem the more pressing issue of disagreement.)

philosophical significance of the perspectivalist's formal semantics lies. Of course, it will not be possible to respond completely to all these worries, but I hope to say enough to indicate why one might not take the task of doing so to be hopeless.

4. The Distinction between Perspectivalism and Realism

With this in mind, let's consider the first of these worries. At first, this worry may appear misplaced. After all, what could be more radical than the thesis that the normative truth is sensitive to the perspective of the assessor? But remember that the notion of truth_S with respect to which this is true for the perspectivalist is a technical one that differs from our everyday, monadic notion of truth_O . Thus, in order to understand the philosophical significance of perspectivalism, we need to understand what is special about this technical notion of truth_S . And, in particular, we need to understand what it is about this notion of truth_S that makes it impossible for the realist to accept the claims the perspectivalist makes using it.

This challenge is particularly pressing because it is plain that the realist can accept claims that look a great deal like the claims the perspectivalist makes in giving his formal semantics. As John Hawthorne and Herman Cappelen note when discussing forms of relativism about matters of taste that are akin to perspectivalism:

[A] realist can perfectly make room for a family of properties expressed by constructions of the form 'true by so-and-so's standards', properties that are distinct from those of truth and falsity. Adopting now the perspective of such a realist, it will be natural to interpret the relativist's talk of some proposition being true at a standard of taste index as expressing the claim that the proposition is true by such and such standards, a perfectly legitimate claim even by the realist's lights. Meanwhile, it will be very natural to interpret the relativist's disquotational truth predicates as expressing the very properties that the realist expresses by 'true' and 'false'. According to this proposed translation manual, the so-called relativist and the realist do not differ at all.³²

In this way, it is *prima facie* easy to establish a translation manual between the perspectivalist's account of normative discourse and that of the realist so that there does not appear to be any real disagreement between them—only an attempt by the perspectivalist to re-label the familiar property of 'being true by so-and-so's standards' in such a way as to generate a position of real philosophical interest. If this is right, then—much like certain forms of sophisticated expressivism—there is

³² In Cappelen and Hawthorne (2009, p. 137).

real cause for concern that there is nothing of substance at issue between the realist and the perspectivalist—in which case the perspectivalist will hardly be in a position to claim that his view represents an advance beyond standard forms of realism.

The perspectivalist can best meet this challenge by telling us something about his understanding of the technical notions of truth_S and falsity_S at work in his semantics that makes it impossible for the realist to translate this semantics into terms he can accept. Thus, this challenge and the challenge to make good sense of the very idea of relative truth are closely related. Along these lines, John MacFarlane—arguably the leading contemporary proponent of this kind of view—responds to this objection as follows:

To see the difference, we need to look at what the realist and the relativist say about accuracy. The realist will say that the accuracy of an assertion that something is tasty is independent of anyone's standard of taste, while the relativist will say that it depends on the standard of taste relevant at the context of assessment. Given plausible principles connecting accuracy to norms for assertion and retraction, the two views will imply different things about which tastiness assertions are warranted, and hence (indirectly) make different predictions about usage.³³

According to MacFarlane, it is possible to make sense of the idea of truth_S relative to (the perspective of) an assessor in a manner that distinguishes this notion of truth_S from anything a realist can accept, because the perspectivalist's notion of truth_S will have normative consequences that are different from the consequences that follow from an attempt to translate perspectivalism into terms the realist can accept. In other words, for MacFarlane, perspectivalism does not collapse into realism because the differences between them *matter* to what we are warranted in doing when we speak and think.

This is a reasonable way of thinking of the issue before us. After all, if there is no normatively significant difference between the realist's account of an area of discourse and the perspectivalist's, then we should be suspicious about whether they are more than mere notational variants. But if we can identify some such difference, this worry will disappear.

Can we, then, identify a pattern of usage that would be warranted given perspectivalism, but not given normative realism? To do so, according to MacFarlane, we need to focus on two aspects of our linguistic practice: how and when we make new assertions and how and when we retract assertions that we have already made. The first of

³³ MacFarlane (forthcoming b, p. 7).

these, according to MacFarlane, is what distinguishes the normative implications of perspectivalism from those of the realism—while the second of these distinguishes the normative implications of perspectivalism from those of a contextualism that treats normative truth as Judgment-Dependent in the manner described above. In short, according to MacFarlane, an assessment-sensitive account will be appropriate for some area of discourse if the following facts hold of that area:³⁴

1. Realism makes assertion and judgment too hard in the area of discourse. If realism is true, then we are all making many assertions and judgments that we are not warranted in making. Thus, a realist account of the area of discourse cannot do justice to our actual practice of assertion and judgment.
2. Contextualism makes retraction too hard in the area of discourse. If contextualism is true, then we should not retract our previous assertions or judgments in the face of a shift in our state of mind as often as we do.³⁵ Thus, a contextualist account of the area of discourse cannot do justice to our actual practice of retraction.³⁶

According to MacFarlane, it is only by conceiving of truths_s in assessment-sensitive terms that we are in a position to capture both of these aspects of certain areas of discourse. Thus, we can see what is special about the perspectivalist's position by focusing on the *norms* that govern an assessment-sensitive area of discourse. In saying this, it is important to stress that one need not endorse any particular account of the *source* of these norms or of the *nature* of assertion. On some accounts, norms of these sorts will be *constitutive* of the acts of judgment, assertion, and retraction within these areas of discourse.³⁷ But this need not be the case—for example, it may be possible to derive

³⁴ Note that this is not meant to be an exhaustive characterization of when such an account might be appropriate.

³⁵ MacFarlane claims that the phenomenon of retraction only makes sense in the case of assertion as opposed to judgment, but this seems to me a mistake. There are norms governing when we are required to revise our judgments that are parallel to the norms governing the retraction of assertions.

³⁶ MacFarlane takes this point to count equally against indexical and non-indexical forms of contextualism, but whether this is the case may depend somewhat on the pragmatics we pair with the non-indexical contextualist's semantics. If it is in fact possible to find a non-indexical contextualist pragmatics that avoids this result, then it may turn out to be the case that the resulting form of non-indexical contextualism is a notational variant of what I am calling perspectivalism. For more discussion of this issue, see my (forthcoming c).

³⁷ For two prominent defenses of the idea that assertion is governed by a constitutive norm of this sort, see Dummett (1959) and Williamson (1996).

these norms from more basic facts about, say, the purpose of these acts or the commitments they involve.³⁸ I will remain agnostic about these issues here—provided, of course, that it is possible to make sense of these acts as governed by linguistic norms of the general sort discussed.

Given this, the perspectivalist can argue that his account is well placed to capture the relevant phenomena as follows. First, in order to avoid making judgment/assertion too difficult, the perspectivalist need only conceive of the norms governing judgment/assertion so that whether one is permitted to make a judgment/assertion depends on whether that judgment/assertion is true_S relative to one's current perspective.³⁹ For example:

Judgment/Assertion-Relative: One is permitted to judge/assert that P just in case P is true_S relative to one's own perspective.⁴⁰

If we conceive of the norms governing judgment and assertion in this manner, there is no difficulty in explaining why it is permissible to form normative judgments in the relatively free manner we normally do.⁴¹ Similarly, in order to avoid making retraction too difficult, the perspectivalist need only conceive of the norms governing retraction so that one is required to make a retraction just in case the judgment/assertion one retracts is not true_S relative to one's current perspective:

Retraction-Relative: One is required to retract one's previous judgment/assertion that P just in case P is not true_S relative to one's own perspective (and the truth of P is currently being challenged).

Thus, the argument goes, the theorist who conceives of truth_S in assessment-sensitive terms is ideally positioned to get both of these

³⁸ For the idea that assertion has an essential purpose or effect see Stalnaker (1978/1999). For the idea that assertions should be understood in terms of commitments, see Brandom (1983).

³⁹ See MacFarlane (forthcoming a, ch. 5). As MacFarlane notes there, the basic point here may also be captured in the terms of Brandom's commitment-theoretic account.

⁴⁰ Of course, this norm can be modified in any of the various ways familiar from the contemporary debate about the norms governing assertion—e.g. so as to involve knowledge or justification as opposed to truth. (There is also the delicate issue here concerning whether speech acts that are governed by this norm should count as 'assertions' in a strict sense. Since this is primarily a question of terminology, I set it aside here.)

⁴¹ Not that making such judgments or assertions should be trivially easy, of course—but, given the degree of idealization that goes into the notion of a normative perspective, there is no danger of this.

dimensions of usage right. By conceiving of the norms governing judgment, assertion, and retraction in terms of assessment-sensitive truth, this theorist claims that he can give an elegant explanation of why we are entitled (or required, as the case may be) to make the judgments, assertions, and retractions we generally do.⁴²

Now, in considering these norms, it is crucial here not to read them as norms of *correctness*—at least if we think of correctness as itself being assessment-independent. For, on the assessor-sensitive view, the fact that an assertion is permissible with respect to the above norms does not mean that this assertion is ‘correct’ in the ordinary sense of this word. After all, given the norms just described, when one’s perspective shifts, one will sometimes be required to retract a previous assertion, even though this assertion was permissible given the norm stated in Judgment/Assertion-Relative. And yet surely it does not make sense to demand of someone that they retract a ‘correct’ assertion—although it certainly can make sense to ask someone to retract an assertion that was perfectly permissible in the context it was made.

For this reason, the perspectivalist should continue to think of the ‘correctness’ of judgments and assertions in terms of the familiar equation between ‘correctness’ and truth_O.⁴³ Similarly, the perspectivalist should not think of these norms as *mere* norms of epistemic warrant or justification. Rather, he should insist that neither the norms of ‘correctness’ nor the norms of epistemic warrant or justification are what is explanatorily fundamental in cases of assessment-sensitive discourse. For in such cases, these norms will both be explained by the more fundamental linguistic or conceptual norms just described, which characterize the correct usage of the terms or concepts at issue.⁴⁴

Here, though, the reader is likely to protest. For how can we appeal to such norms in making sense of the perspectivalist’s proposal when this proposal is *itself* a proposal about how to understand normative semantics? After all, if the norms we are appealing to are themselves subject to an assessor-sensitive account, then in making reference to

⁴² Given that the ‘cash value’ of the notion of truth_S relative to a perspective is primarily a matter of norms of this sort, one might wonder whether it would be better to leave all talk of truth-conditions aside and simply speak of a formal notion of, say, ‘appropriateness’ that underlies these norms. I leave this debate to the side here.

⁴³ Or perhaps, *à la* Williamson, in terms of *knowledge* of truth_O.

⁴⁴ The possibility of such norms should not, I think, be terribly mysterious—since many areas of discourse are governed by systems of norms that cannot be accurately captured either in terms of ordinary notions of ‘accuracy’ or ‘correctness’ or in straightforwardly epistemic terms. The perspectivalist’s claim, in effect, is simply that the fundamental linguistic norms governing normative discourse have this character.

these norms, we have made no real progress in making sense of assessment-relative truth. But if these norms are not subject to such treatment, what's the point of insisting upon an assessment-dependent account of *any* norms at all?

Fortunately, the linguistic norms we are discussing now are quite different from most norms in ways that make giving an absolutist account of these norms *relatively* easy. In particular, these norms do not necessarily raise the sorts of metaphysical or epistemological issues that motivate constructivism. For example, they might reasonably be regarded as the product of linguistic or conceptual conventions. Of course, even if these norms are conventional in some sense, there is the further question of whether we ought to accept linguistic conventions of the relevant sort. And this further question will be the sort of substantive normative question that warrants perspective-dependent treatment for the perspectivist. But this is perfectly compatible with the role these norms are playing in the present account.

With this in mind, let's return to the main topic of this section. Since I am concerned here with the possible differences between normative realism and normative perspectivalism, it is the first of the two potential differences noted above that is important here. The crucial question, then, is whether normative realism really does make it 'too difficult' to make normative judgments and assertions. Here the normative realist is sure to protest that it is perfectly possible to give a plausible normative epistemology that is compatible with realism *and* that vindicates the relative ease with which we generally make normative judgments and assertions. To assume otherwise, he will insist, is simply to treat the realist as if his view forces us into normative skepticism. And this is yet to be shown.

Of course, there are arguments that purport to show just this. But whether or not these arguments are successful is something we need not consider. For our interest in constructivism was sparked by the worry that these arguments (or certain closely associated metaphysical arguments) *might be* successful. Thus, when evaluating the question of whether perspectivalism is merely a notational variant of realism, we may *assume* that it has, for some reason, proved impossible to give a satisfactory realist epistemology and metaphysics in the normative case. For only in this case will we have more than a purely academic interest in perspectivalism.

As such, there are two possibilities here. First, it may be possible to develop a normative epistemology that vindicates our current practice of normative assertion and judgment against the background of normative realism. In this case, it is unclear whether there would be any normatively significant difference between a perspectivist

account of normative discourse and a realist one. But in such a case, this need not trouble us—for if this is true, we have no need of a perspectivalist account, given the (assumed) vindication of normative realism.

On the other hand, it may prove impossible to develop a normative epistemology that vindicates our ordinary practices in this area against the background of normative realism. In this case, we will have good reason to look for a view other than normative realism. And fortunately, in this case, the manner in which the perspectivalist conceives of the norms governing normative talk and thought will generate different conclusions about whether our practices are warranted than the realist's account does.

In particular, the normative realist will understand the norms governing warranted judgment and assertion to demand (say) sufficient evidence or knowledge of objective normative facts, as in the following examples:

Evidence: The judgment/assertion that *P* is warranted just in case the judge/asserter has sufficient evidence in favor of *P*.

Knowledge: The judgment/assertion that *P* is warranted just in case the judge/asserter knows that *P*.

It is these demands, of course, that generate the familiar epistemological worries about the realist's position. But, for the reasons noted above, the perspectivalist will not understand the norms governing warranted judgment and assertion in this manner. On the contrary, the fundamental norm governing judgment/assertion for the perspectivalist will have the following form:

Judgment/Assertion-Relative: One is permitted to judge/assert that *P* just in case *P* is true_S relative to one's own perspective.

Thus, for the perspectivalist, if we want to consider when someone is warranted in making the judgment/assertion that *P*, we need to consider her epistemic position with respect to the question of whether *P* is true_S relative to her normative perspective. In other words, we should conceive of the norms governing warranted judgment/assertion in something like the following fashion:

Evidence-Relative: The judgment/assertion that *P* is warranted just in case the judge/asserter has sufficient evidence in favor of *P* being true_S relative to his normative perspective.

Knowledge-Relative: The judgment/assertion that P is warranted just in case the judge/asserter is a position to know that P is true_S relative to his normative perspective.⁴⁵

And it is not difficult to explain how we might generally be in a position to satisfy these norms. In this way, the perspectivalist can explain why normative judgments and assertions should be relatively easy to make in a warranted fashion without taking on the burdens involved in providing a non-skeptical realist epistemology. And so, if it becomes apparent that the project of providing such an epistemology is doomed, the perspectivalist is well placed to account for our practice of normative judgment and assertion by other means.⁴⁶

In this way, we can answer the first challenge to perspectivalism disjunctively. Either there is no real difference between realism and perspectivalism, in which case realism is all we need. Or there is a real difference between the two views, which is fortunate because we need a real alternative to realism. Thus, perhaps surprisingly, perspectivalism emerges as a distinctive view of normative talk and thought just in case we require a view like it. Whether or not such a view is required is, to my mind, very much an open question. But if it is required, then there need not any concern about whether perspectivalism is simply realism in disguise.⁴⁷

5. The Intra- and Interpersonal Problems of Disagreement

In this way, the task of distinguishing perspectivalism from realism is not nearly as hopeless as it might seem, on the supposition that an alternative to realism is required. With this in mind, let's turn to the second main worry about the perspectivalist's account—the issue of disagreement.

⁴⁵ Or, alternatively, to avoid issues involving possession of the concept of truths: the judgment/assertion that P is warranted just in case it is true_S relative to the judge/asserter's normative perspective that the judge/asserter knows that P.

⁴⁶ Note that the claim that the norms governing when a judgment or assertion is *permitted* or *warranted* have this character does not *obviously* conflict with our ordinary judgments about these issues.

⁴⁷ It would be a mistake to think that the perspectivalist should *also* accept the norms of Evidence and Knowledge that the realist accepts. For example, with respect to Knowledge, the mere fact that P is false_S relative to the perspectivalist's context of assessment is sufficient to make it false_S that S knows that P relative to his context of assessment as well. But this fact alone should not necessarily be taken by the perspectivalist to make S's judgment that P unwarranted. Thus, the perspectivalist may regard Evidence and Knowledge as overly demanding variants of Evidence-Relative and Knowledge-Relative that the realist has come to accept only because of his mistaken views about normative discourse.

MacFarlane also has a great deal of use to say about these worries. But for reasons that will become clear shortly, what he says seems to me to leave us short of a fully satisfying response to them. Nonetheless, with some modifications, I believe that a view like his provides us with the elements needed to construct a reasonably plausible account of the sort of disagreement that is present in normative discourse.

MacFarlane's official account of the sort of disagreement that is present in the assessment-sensitive case begins by distinguishing several different forms of disagreement and exploring which of these apply to assessment-sensitive discourse. To see what these different forms of disagreement involve, suppose we have before us two individuals, one of whom judges that *P* and one of whom judges that not *P*. And suppose that the truth_S-value of *P* varies among contexts of assessment such that both individuals' judgments are true relative to their own contexts of assessment. Finally, suppose that all of this is common knowledge between them. In what sense, according to MacFarlane, can these two individuals be said to disagree with one another?

MacFarlane argues that there are at least two important forms of disagreement present in such cases.⁴⁸ First, given the norms for judgment and assertion noted above, it will be impossible for any person to rationally accept both of the judgments in question. For, putting aside issues of vagueness and indeterminacy, if one of these judgments is true relative to a context of assessment, the other will be false. Moreover, this is a trivial conceptual truth about judgments of this form that will be obvious to any one who grasps the concepts involved in them. Thus, it will be obvious that it is impossible to comply with the norms governing these judgments while accepting both of them.

In this way, it is impossible rationally to accept both of these judgments at a single time. Or, as MacFarlane puts the point, these judgments are *doxastically non-cotenable*. But while this may represent a form of disagreement, as MacFarlane himself stresses, *on its own* it is a rather weak form thereof.

To see why, consider a theory of tense according to which propositions are themselves tensed. Given such a theory, my utterance on Monday of 'it is raining now' will express the same proposition as my utterance on Tuesday of this phrase. Thus, given this theory, the proposition expressed by my utterance on Monday of 'it is raining now' will *not* be rationally cotenable with the proposition expressed by my utterance on Tuesday of 'it is *not* raining now'. And yet this obviously does not indicate any real disagreement between my Monday self and my

⁴⁸ For the most accessible presentation of MacFarlane's taxonomy of different forms of disagreement, see his (forthcoming c).

Tuesday self about when it is raining. Thus, unless we are willing to rule views of this sort out of consideration on these grounds—as authors like Hawthorne and Cappelen do—there will be cases in which mere doxastic non-cotenability does not indicate the presence of the sort of ‘genuine’ disagreement we normally take to be possible in the normative domain.

MacFarlane is well aware of this fact, so he stresses that there is a further form of disagreement present in cases of assessment-sensitivity. In particular, as he notes, the doxastic non-cotenability present in such cases is a product of one of their further features namely, that it is impossible for both these judgments to be true_S relative to a single context of assessment. Or, as MacFarlane puts, these judgments display *preclusion of joint accuracy*.

Unfortunately, as MacFarlane himself is well aware, noting this fact *on its own* does little to establish that there is an interesting further form of disagreement present in these cases. After all, as we have stressed repeatedly, truth_S is a technical notion that lacks obvious philosophical or normative significance on its own. Thus, the fact that two judgments cannot both be true_S relative to a single context of assessment does not necessarily show that they are in conflict with one another in any intuitive or philosophically interesting sense. To show that, we need to demonstrate that this sort of preclusion of joint accuracy carries with it the kinds of normative consequences that we associate with genuine disagreement. And to do this, we must go beyond the taxonomy of different forms of disagreement that MacFarlane provides.

Fortunately, in his account of the norms that govern assessment-sensitive discourse, MacFarlane provides us with some of the materials necessary to accomplish this task. But when we turn to these norms, we discover an interesting feature of MacFarlane’s view: it treats intra-personal disagreement with one’s past self differently from interpersonal disagreements with others.

In particular, MacFarlane is able to give a reasonably attractive account of the normative significance of the preclusion of joint accuracy in the intrapersonal case because of his norm for retraction. Given Retraction-Relative, one will be required to retract any (relevant) previous judgment or assertion if it is not true_S relative to one’s current perspective. And when one retracts a previous judgment or assertion, there is an obvious sense in which—by doing so—one is disagreeing with this judgment or assertion. For the ordinary speech act of retraction is one of the acts that is characteristic of the state of disagreeing with one’s past self.

Thus, we can say that whenever one is in a state of mind that makes it the case that one of one’s previous judgments or assertions is not

true, relative to one's current perspective, it is reasonable to say that one disagrees with this judgment or assertion precisely *because* one is in a state of mind that requires one to retract it. In this way, the perspectivist can account for the possibility of intrapersonal disagreement by beginning with an ordinary *act of disagreement*—namely, the act of retraction—and then considering the conditions under which this act of disagreement is appropriate vis-à-vis one's previous judgments or assertions.⁴⁹ By doing so, he can explain, at least in the intrapersonal case, why the preclusion of joint accuracy represents a genuine form of disagreement by pointing to the normative consequences of this state. In effect, by following this line of thought, he can argue that the intrapersonal preclusion of joint accuracy deserves to be considered a genuine form of disagreement because it requires one to perform actions that are characteristic of the state of disagreeing with one's past self.

An account of disagreement of this general sort has a good deal to recommend it. But while an account of this sort falls naturally out of MacFarlane's account of the norms governing assessment-sensitive discourse in the intra-personal case, the same is not true in the interpersonal case.

In a certain sense, this should be no surprise—for MacFarlane does not provide an interpersonal counterpart to the retraction norm which gives real normative bite to disagreement *qua* preclusion of joint accuracy in the intrapersonal case. As a result, when we consider a case of interpersonal preclusion of joint accuracy like the one described above, there is nothing in MacFarlane's norms that requires me to *actively disagree* with those whom I 'disagree with' in this sense. To be sure, when I encounter someone who believes something which conflicts with my beliefs in this way, I am prohibited by MacFarlane's norms from coming to share this belief without changing my other beliefs. But there is nothing in these norms that implies that I should try to argue with, refute, or reject these beliefs or claims. Nor is there anything in these norms that suggests that I should feel any rational pressure to bring it about that our beliefs no longer conflict in this way. Rather, so long as I do not shift my own point of view, we can both adopt a 'live and let live' attitude about this difference in our beliefs and go on our merry way.⁵⁰

This having been said, there is still an important difference between cases of interpersonal doxastic non-cotenability and cases of interpersonal preclusion of joint accuracy as described by MacFarlane. For in the latter case, but not the former, if I somehow convince some-

⁴⁹ This, I should note, is not exactly the account of disagreement that MacFarlane provides— but it strikes me as the most attractive way of using the resources available to him.

⁵⁰ Similar issues arise with respect to MacFarlane's assessment-sensitive version of the Brandomian commitment-theoretic account.

one to change her mind about the issue in question, she will be required to *retract* her previous views about it. And this does seem to capture a more robust sort of disagreement than mere doxastic non-cotenability on its own. But nonetheless it does nothing to explain why I should feel the need to try to bring it about that those who differ in their views retract these views in this way. Nor does it explain why my disagreement with them should generate any rational pressure on them to do so. Rather, any pressure they feel to shift their views in the face of such disagreement will be the product of contingent features of their psychology—such as the contingent displeasure associated with not sharing someone’s views—as opposed to anything about the norms governing this area of discourse.⁵¹ And if this is the only sort of disagreement that is present in cases of inter-personal preclusion of joint accuracy involving assessment-sensitive terms, does this sort of preclusion of joint accuracy *really* deserve to be regarded as a form of ‘genuine’ disagreement at all?

In particular, the sort of interpersonal ‘disagreement’ that we have identified thus far seems weaker than the disagreement that we find in normative discourse. For when someone disagrees with me about a normative matter, I often take this to put *some* rational pressure on me to either refute his view or explain it away. And surely we do generally feel compelled to engage in normative disputes to a degree the above discussion does not explain. Thus, if we hope to give a perspectivist account of normative discourse, we will need to give a more compelling account of interpersonal disagreement than we have thus far.

Moreover, these issues are symptomatic of a more basic issue with the account of assertion that we have been considering thus far: namely, that it fails to do justice to the basic Stalnakerian insight that one (if not the) main purpose of the speech act of assertion is to generate convergence within a conversational context via adding to the set of propositions that are part of the shared common ground within that context.⁵² The norms we have been discussing so far fail to do justice to this idea because, when we consider an assessment-sensitive assertion, there is nothing in these norms that captures the idea that the aim of assertion is to increase a conversation’s common ground in this way. For, as we have already

⁵¹ This feature of MacFarlane’s views comes out towards the end of his (2007). We might try to supplement this account by appealing to our contingent interest in achieving the relevant forms of convergence, but this would still raise the issues of concern below.

⁵² See, in particular, Stalnaker (1978/1999). For a quite different attempt to think through Stalnakerian ideas in an assessment-sensitive context, see Egan (2007). MacFarlane does discuss Stalnaker’s account in his (2011), but he does not try to make sense of relative truth in terms of an account of this sort.

noted, in the assessment-sensitive case there is nothing in these norms that puts the members of a conversation under any pressure to arrive at a shared view of the matter they are discussing. And so there is nothing about these norms that captures the thought that an assertion that does not produce this result has been *less than successful*.

In order to capture this feature of interpersonal normative disagreement, we need to add a further element to the norms we have discussed thus far. Just as in the intra-personal case, this additional element should capture the appropriateness of responding to cases of interpersonal preclusion of joint accuracy with those speech acts that are characteristic of genuine inter-personal disagreement. And it should also do justice to the idea that assertion, as such, aims to produce convergence on a shared conversational common ground.

Fortunately, if what we have said is correct, the picture sketched above of intrapersonal disagreement provides us with a natural model for how we might accomplish these tasks. In particular, in order to do so, we must simply add a further norm to our account that captures the actions that are characteristic of *interpersonal* disagreement.

To do so, we must find some act or acts of interpersonal disagreement that can play the role that retraction played in our account of intrapersonal disagreement. Here it is important that the act or acts we choose should be familiar from our *ordinary* experience of active disagreement. For only then will they connect the formal apparatus of assessment-sensitive truths-conditions together with our ordinary understanding of what it is to disagree. Moreover, the acts in question should not be acts that are made appropriate simply by *a desire to change the mind* of the person one is in conversation with. For, of course, one can desire to do this in cases in which no disagreement is present. Rather, what we are looking for are speech acts that are characteristic of the sort of disagreement that is present in normative discourse. Fortunately, there are certain speech acts that are well placed to play this role—speech acts that (at least in normal conditions) are appropriate if and only if two parties disagree with one another in the sense of interest to us here. And this could hardly fail to be the case—for if there were no speech acts that were distinctive of the sort of ‘genuine’ disagreement that is present in normative discourse, it is hard to know how we could tell that this sort of disagreement is present in the first place.

One possible speech act we might point to as playing this role is the speech act of *rejecting* what someone else has said. For it does seem right to say that we should only reject someone else’s assertion when the sort of ‘genuine’ disagreement we are interested in is present. In this way, the speech act of rejection might be taken to play an analogous role in interpersonal disagreement to the role that retraction plays in

the intrapersonal case.⁵³ But one might deny that there really is a speech act of rejection with these features. Fortunately we can sidestep this issue by focusing on another phenomenon that is distinctive of interpersonal disagreement of the sort we are interested in here—namely, the phenomenon of *attempting to refute* someone else's view—or, in other words, of arguing against that view with the aim of *bringing it about that he retracts his former view*. This activity too appears to be distinctive of the sort of disagreement we are interested in. And by focusing on the appropriateness of this activity, we can get a grip on the nature of interpersonal disagreement in a manner that builds upon our previous understanding of the role of retraction in the intrapersonal case without relying on a new speech act of rejection.

Taking up this idea, we might say that it is only appropriate to *dispute* someone's assertion if one genuinely disagrees with it. If this is correct, then we can explain why the interpersonal preclusion of joint accuracy represents a robust form of disagreement, even in the assessment-sensitive case, by connecting together these speech acts with truth_S-conditions in something like the following manner:

Dispute-Relative: If a party to one's conversation has made a straightforward assertion that P and P is not true_S relative to one's own perspective, then (provided P remains relevant) one is required to try to get that party to retract this assertion so long as one remains in conversation with him. Until one does so, one cannot regard that conversation as fully successful.⁵⁴

In reading Dispute-Relative, it is important to remember that it does not require one to dispute a normative point endlessly with those who do not share one's point of view. Rather, all that it requires is that one dispute any such point for as long as one is in a conversation about it with those who do not share one's views. Thus, these norms force one to make a choice: either dispute the relevant differences of opinion or give up on the project of successfully conducting conversation with those who disagree with one. In this way, they nicely capture the Stal-

⁵³ Once again, it is important to stress that my talk of 'rejection' is meant to pick out an ordinary speech act with which we are all familiar. In this respect, the appeal to rejection here is very different from the appeal to related notions one finds in, say, the work of Allan Gibbard—which to some degree do appear to be technical notions introduced for the purposes of theorizing about the meaning of normative terms and concepts.

⁵⁴ The reference to a 'straightforward assertion' is meant to rule out assertions that take place in conversational contexts in which it is possible to accept a proposition for the purpose of that conversation without actually coming to believe it. Hopefully it is clear how this principle could be extended to deal with such cases.

Stalnakerian idea that the purpose of successful assertion and conversation is convergence on a shared view of the issues under discussion.

In particular, given Dispute-Relative, there are only two ways such a conversation can end. First, one or both parties can retract their previous assertion so that they no longer have divergent views on the question at issue. Or, second, one can give up on the project of conversing about the question at hand—in effect, taking the other party to be unworthy of engaging with about this issue. Crucially, the choice between these two options provides a further source of rational pressure towards convergence under normal circumstances. For if we wish to avoid simply writing off those we are in conversation with as not worth talking to, we can reach a stable and successful conclusion to such a conversation only via bringing it about that at least one party retracts his previous view of the matter. And, given the norms for assertion and retraction above, it will only be permissible for the parties to the conversation to retract their previous views insofar as their normative perspectives have themselves shifted in corresponding ways. Thus, we can only reach a stable and successful conclusion to a conversation of this sort by achieving convergence in our normative perspectives—at least insofar as they are relevant to the issues under discussion.⁵⁵ Or:

Success: If I have made a straightforward assertion that P, then I am permitted to regard this assertion as fully successful only when P is true_S relative to the perspective of every party to the conversation.⁵⁶

At the same time, this source of rational pressure is not so strong that it renders moot our attempts to make permissible judgment and assertion relatively easy. And yet, Dispute-Relative allows us to explain why interpersonal preclusion of joint accuracy should count as a ‘more genuine’ form of disagreement by connecting it with our ordinary understanding of what such disagreement involves. In doing so, it is true that we begin with the acts of disagreement that we are all familiar with from our ordinary lives and use the fact that a certain state makes these acts appropriate to justify the claim that this state represents a genuine form of disagreement. But it is difficult to see how else we should proceed here—

⁵⁵ Of course, an exchange of assertions can be partially successful even when this fails to occur. But no assertion that does not have this effect can be regarded as wholly successful.

⁵⁶ Again some of these norms may be derivable from more fundamental norms or aims, as is suggested by some of Stalnaker’s discussion of these issues. For example, we might see many of these norms as the product of the basic idea that normally, in asserting P, I am trying to get you (i) to judge that P and (ii) to retract any previous judgments or assertions that not P.

given that our pre-theoretical grasp of what it is to disagree with someone is presumably largely constituted by our understanding of the speech acts that such disagreement makes appropriate.

Moreover, insofar as we have a pre-theoretical understanding of disagreement in *semantic* terms, this understanding must involve the ordinary concepts of truth_O and falsity_O—and these concepts do nothing to distinguish the perspectivist's account of disagreement from the realist's. Thus, if we are to get any useful purchase on our ordinary understanding of the nature of disagreement here, it seems that this must come from our understanding of the speech acts that disagreement makes appropriate.

Nonetheless, in proceeding in this fashion, it may be thought that we have yet to *really* explain why the states in question represent a genuine form of disagreement. For at most we have claimed that the basic norms governing normative discourse make it appropriate to respond to cases of inter-personal preclusion of joint accuracy with the speech acts that are characteristic of 'genuine' disagreement. And this, it may well be felt, does not really explain why these norms make sense. Rather, it simply involves the *stipulation* that they do. After all, one may think, it makes sense to respond to the preclusion of joint accuracy in this way in the non-assessment-sensitive case *because*, in this case, this state indicates a dispute about some non-relativistic matter of fact. And, of course, no such explanation can be provided for the appropriateness of these speech acts in the assessment-sensitive case.

6. The Point of Perspectivalism

This raises a real question about the norms that the perspectivist appeals to in making sense of assessment-sensitivity. For surely it is incumbent upon the perspectivist to provide an account of these norms that makes it plain *why* it would make sense to have a linguistic practice with the features he describes. But this is less a question about the perspectivist's ability to capture a sufficiently robust form of disagreement within the confines of his theory and more a question about what the *point* of the sort of linguistic practice he describes would be. That is, the question here is less whether a linguistic practice involving these norms would be possible—for it is hard to see what would rule it out *in principle*—and more whether it would ever make sense to introduce concepts and terms into our language that obeyed the sorts of norms the perspectivist appeals to.

To this the perspectivist should reply that the sort of discourse he has described makes sense just in case it makes sense to have a linguistic practice that aims at producing convergence in our normative perspectives through conversation without thereby committing us to normative

realism. And, once again, if we are skeptical about normative realism, the utility of such a form of linguistic practice seems hard to deny.

To flesh out this point, the perspectivalist might return to the remarks we made above in our attempt to distinguish perspectivalism from realism. For what, after all, do we want our practice of normative discourse to accomplish? Putting aside the practicality of normative judgment, we want our normative discourse to have at least two main features. First, it should be relatively easy to form normative judgments and to make normative assertions on the basis of one's own normative perspective. And, nonetheless, part of the character of normative assertions should be that they aim at bringing it about that those one is talking to retract views that differ from one's own. This is important insofar as we are interested in using normative discourse as a mechanism for achieving *coordination* in our normative perspectives. For an area of discourse that supports a norm like Dispute-Relative is far better placed to fulfill this role than an area of discourse that does not.⁵⁷

Once again, if the perspectivalist's pessimism about normative realism is correct, the normative realist is poorly placed to arrive at a view with the first of these features. And yet views that involve, say, the Subject-Dependence or the Judge-Dependence of Normative Truth do not seem to support the acceptance of a norm like Dispute-Relative, and so do not generate the same sort of pressure to convergence that the perspectivalist's does. Thus, if the perspectivalist's pessimism about normative realism is warranted, the sort of linguistic practice the perspectivalist describes appears to be especially well suited to fulfill these aims.

In this way, perspectivalism is excellently placed to serve as an alternative to normative realism, at least for those who are sympathetic to the constructivist's basic conviction that the truth of normative judgments is perspective-dependent in *some* way. Given this, it is worth stressing once again the similarities between the perspectivalist's position and the views of sophisticated contemporary expressivists who share similar motivations. But while there are deep similarities between these views, we should be cautious about too quickly identifying them. For there is at least this *apparent* difference between them: contemporary expressivists arrive at their account of the meaning of normative judgments via an attempt to complicate traditional forms of expressivism so as to respond to familiar objections to these views. By doing so, they inevitably reproduce much of the structure of traditional truth-conditional semantics. But the significance of this structure remains rooted, at least to some degree, in the ambitions of the traditional expressivist pro-

⁵⁷ MacFarlane himself, of course, stresses the importance of coordination here, as do Gibbard (2003) and Price (2003).

ject. The perspectivalist, as I understand him, approaches these issues from the opposite direction. He begins with traditional truth-conditional approaches to normative semantics—be they realist or constructivist—and then attempts to modify these approaches to capture the sort of perspective-dependence he claims to find in normative discourse. Of course, it is possible—and perhaps even likely—that end result of the expressivist’s journey away from traditional forms of expressivism and the end result of the perspectivalist’s journey away from traditional forms of constructivism will converge. But we should not assume that this will transpire—and we should certainly not assume that this convergence has already occurred.⁵⁸

References

- Armstrong, David. 1997. *A World of States of Affairs*. Cambridge: Cambridge University Press.
- Blackburn, Simon. 1998. *Ruling Passions*. Oxford: Oxford University Press.
- Brandom, Robert. 1983. “Asserting.” *Noûs* 17: 637–650.
- Brogaard, Berit. 2008. “Moral Contextualism and Moral Relativism.” *Philosophical Quarterly* 58: 385–409.
- Cappelen, Herman and Hawthorne, John. 2009. *Relativism and Monadic Truth*. Oxford: Oxford University Press.
- Dreier, James. 1990. “Internalism and Speaker Relativism.” *Ethics* 101: 6–26.
- . 2004. “Metaethics and the Problem of Creeping Minimalism.” *Philosophical Perspectives* 18: 23–44.
- . 2009. “Relativism (and Expressivism) and the Problem of Disagreement.” *Philosophical Perspectives* 23: 79–109.
- Dummett, Michael. 1959. “Truth.” *Proceedings of the Aristotelian Society* 59: 141–162.
- Egan, Andy. 2007. “Epistemic Modals, Relativism, and Assertion.” *Philosophical Studies* 113: 1–22.
- Enoch, David. 2006. “Agency, Schmagency: Why Normativity Won’t Come From What is Constitutive of Action.” *Philosophical Review* 115: 169–198.

⁵⁸ This having been said, there is some debate in the contemporary literature about how positions like Gibbard’s are best understood. So the reader may take the view that the perspectivalist’s account is, in the end, best thought of as a way of re-characterizing Gibbard’s account in more familiar truth-conditional terms. Even if this is the case, I hope that this re-characterization is a helpful one, insofar as it makes plain the manner in which the resulting view can be seen to draw upon the same resources as standard truth-conditional accounts. This alone, it seems to me, would represent a considerable advantage to thinking about these issues in the manner the perspectivalist recommends.

- Gibbard, Allan. 2003. *Thinking How to Live*. Cambridge: Harvard University Press, 2003.
- . Forthcoming. *Thoughts and Oughts*.
- Kastafanas, Paul. Forthcoming. *Agency and the Foundations of Ethics: Nietzschean Constitutivism*.
- Kölbel, Max. 2002. *Truth without Objectivity*. London: Routledge.
- Korsgaard, Christine. 1996. *The Sources of Normativity*. Cambridge: Cambridge University Press.
- Lewis, David. 2001. "Truthmaking and Difference-Making." *Noûs* 35(2001): 602–615.
- MacFarlane, John. 2005. "Making Sense of Relative Truth." *Proceedings of the Aristotelian Society* 105: 321–339.
- . 2007. "Relativism and Disagreement." *Philosophical Studies* 132: 17–31.
- . 2009. "Non-Indexical Contextualism." *Synthese* 166: 231–250.
- . 2011. "What Is Assertion?" In Brown, Jessica and Cappelen, Herman, (eds.), *Assertion*. Oxford: Oxford University Press.
- . Forthcoming a. *Assessment-Sensitivity: Relative Truth and its Applications*.
- . Forthcoming b. "Simplicity Made Difficult."
- . Forthcoming c. "Varieties of Disagreement."
- Melia, Joseph. 2005. "Truthmaking without Truthmakers." In Beebe, Helen and Dodd, Julian (eds.), *Truthmakers: The Contemporary Debate*. Oxford: Oxford University Press.
- Price, Huw. 2003. "Truth as Convenient Friction." *Journal of Philosophy* 100: 167–190.
- Schafer, Karl. Forthcoming a. "Moderate Kantian Constructivism."
- . Forthcoming b. "Quasi-Realism, Projectivism, and the Explanatory Challenge."
- . Forthcoming c. "Retraction, Non-Indexical Contextualism, and Assessment-Relativism."
- Schroeder, Mark. 2007. *Slaves of the Passions*. Oxford: Oxford University Press.
- Setiya, Kieran. 2007. *Reasons without Rationalism*. Princeton: Princeton University Press.
- Sider, Theodore. 2011. *Writing the Book of the World*. Oxford: Oxford University Press.
- Stalnaker, Robert. 1978/1999. "Assertion." In *Context and Content*. Oxford: Oxford University Press.
- Street, Sharon. 2008. "Constructivism about Reasons." *Oxford Studies in Metaethics* 3: 207–246.
- Williamson, Timothy. 1996. "Knowing and Asserting." *Philosophical Review* 105: 489–523.