Ridgeview Publishing Company

ETHICS NATURALIZED

Michael Slote
University of Maryland at College Park

Philosophers who have discussed or advocated naturalized epistemology have often called attention to analogies between ethics or moral theory, on the one hand, and epistemology or theory of knowledge, on the other. And utilitarianism in particular has been singled out as an ethical analogue of a naturalistic or naturalizing approach to questions of epistemic justification.[1] Both BonJour and Firth make the latter point, and since both also regard utilitarianism as ethically unacceptable, they make use of the epistemic/ethical parallelism to cast doubt on naturalizing epistemology.

It is certainly possible, however, to doubt that utilitarianism can readily be dismissed simply on the basis of counterintuitive implications of the sort Firth and BonJour allude to; for it has (for one thing) become increasingly recognized that our intuitive moral thinking can lead to counterintuitive results quite on its own—we have reason to believe, in other words, that common-sense moral intuitions conflict with one another, are incoherent as a class.[2] And this seems to give new life to the philosophical impulse toward theory and system that has been so clearly exemplified in utilitarian ethics. But quite apart from the merits of utilitarianism or naturalistic epistemology, recent discussions connecting these two may give a false impression of that connection by seeming to imply that any epistemological naturalist will inevitably want to adopt some form of utilitarianism if she seeks a coherent overall philosophical view. Utilitarianism naturally lends support to and gains support from naturalized epistemology. But utilitarianism is not the only ethical view that naturally "goes" with a naturalizing approach in epistemology, and in this paper my main purpose will be to point out another, less familiar ethical approach that goes well with naturalized epistemology and show how differences between the less familiar approach and utilitarianism correspond to interesting differences in the ways we can seek to naturalize epistemology. It would not be correct to claim that virtue ethics in general harmonizes with a naturalizing approach to epistemology: quite the contrary, there are stan-

dard forms of virtue ethics that seem most plausibly associated with traditional, internalist epistemology. But there is another, less familiar kind of virtue ethics that can very plausibly be seen as analogous to epistemological naturalism in one of its guises, one whose present-day relevance and importance is further underscored by its ability to deal in a coherent fashion with the exasperating problem of moral/ethical luck. But Utilitarianism also has an interesting way of dealing with this problem, and I propose to begin our discussion by saying something about the paradoxes involved in our ordinary thinking about moral luck.

1. Moral Luck, Kantian Ethics and Reductionist Utilitarianism

One of the unnerving accomplishments of recent discussions of moral luck is to point to a nest of inconsistencies at the heart of our intuitive moral thinking. Consider, for example, our everyday moral reactions to the unforeseen consequences of people's actions. Ordinary moral thinking distinguishes the moral status of attempted murder from that of murder, for example, and quite apart from legal distinctions between murder and attempted murder, we think much worse of someone who has killed an innocent victim than of someone who accidentally fails to kill an intended innocent victim. And we are inclined to heap greater blame on an actual murderer than on an unsuccessful one.

To that extent common-sense morality allows actual unforeseeable consequences a role in determining moral judgment, thus making room for a certain kind of "moral luck". But as has been frequently pointed out, the very idea of moral luck affronts our common-sense moral intuitions. So our moral intuitions about cases taken singly are in conflict with a general common-sense conviction that judgments of morally better and worse, or greater or less culpability or blameworthiness, cannot properly be subject to luck or accident. And the cases where such inconsistency arises are quite numerous and varied.

To take an even clearer example, imagine someone driving a car along a country road and pointing out noteworthy sights to his passengers. As a result of his preoccupation, the car suddenly swerves to the middle of the road; fortunately there are no cars coming in the opposite direction and no accident occurs. However, in another scenario the person is similarly preoccupied, and because a truck happens to be coming, has a major accident. He is then responsible for a great deal of harm to others and would normally be accounted blameworthy or culpable in a way that he would not be thought blameworthy or culpable in the first-mentioned case. This example is borrowed from Thomas Nagel's paper "Moral Luck". But as Nagel also points out, something in us revolts against the idea of moral luck, inclining us to

the view that the driver must have the *same* degree of culpability in the two cases mentioned above.[3] There is something repelling in the idea that one can be more or less culpable depending on events outside one's ken or control. And in regard to the just-mentioned pair of examples, therefore, it may seem as if we should be able to insert some kind of probability estimate into the example, so that whether the driver who swerves is culpable and the degree, if any, of his culpability would depend solely on whether he was sufficiently aware of the likelihood of an accident and on how likely an accident was, given his preoccupation with the scenery—judgments that are constant between the two imagined cases and that might allow us to say the man was culpable (or not culpable) in both cases to the same extent for having paid attention to the scenery while driving.

But (following Nagel) I think that no such solution really squares with the moral judgments we make in the ordinary course of events, before we begin to worry about moral luck in a self-conscious way. I think no matter how constant one imagines the (awareness of) probability in the two situations, common-sense morality sees some difference in the culpability, blameworthiness, or reprehensibleness of the agent. (Actually, it doesn't matter whether the two situations are viewed as counterfactually possible relative to one another or are viewed as involving similar persons—or one person at different times—facing relevantly similar situations.) Common sense appears to subscribe to a mutually contradictory set of assumptions with regard to putative cases of inattention, carelessness, or negligence, and we therefore need to go beyond our ordinary moral thinking.

We have just spoken of one particular form of moral luck: luck with respect to results or consequences of one's actions or inactions. But moral luck can also come about in other ways. Thus if some of us would have done nothing to stop the Nazis, had we been living in Germany before and during the Second World War, then arguably any lesser or negligible guilt we have through not having in fact been in Germany during that period is attributable to moral luck in our *circumstances*. And such assumptions of circumstantial luck lead to inconsistencies that resemble those uncovered in connection with luck in *results* or *consequences*, though I don't want at this point to dwell on these difficulties.

The issue of blameworthiness (or culpability or reprehensibility) is at the very heart of the issue of moral luck, because it is the idea that luck or accident can make a difference to blameworthiness, etc., that most grates against our antecedent moral intuitions. If we concentrate on praiseworthiness the clash of intuitions is less evident, because there is such a thing as non-moral praiseworthiness—we can praise an artistic performance or work that it would make no sense to regard as culpable or blameworthy—and because it is therefore not odd at all to suppose that non-moral praiseworthiness can sometimes depend on accident. Of course, we could distinguish moral praise-

worthiness from praiseworthiness in general and claim that it grates on our intuitions to suppose that moral praiseworthiness can be subject to luck; but it is just easier to focus on blameworthiness, where ambiguity seems less capable of misleading us because all blameworthiness seems to be moral blameworthiness. And so in what follows I shall frame the issues of moral luck largely in terms of the notion of blameworthiness.

The two dominant ethical traditions of recent times—Kantianism and utilitarianism—have different and indeed distinctive ways of avoiding the paradox and inconsistency our intuitive thinking gets into in the area of moral luck. (That is not to say that all, or even most, utilitarians and Kantians have been *self-consciously aware* of the problem of moral luck or that they have adopted their main views in *response* to that problem.) It is fairly clear that Kant wishes to preclude the possibility of moral luck and thus to avoid those judgments about individual cases that jointly imply that someone's (degree of) blameworthiness can depend on factors of luck or accident. The utilitarians, by contrast, make room for the phenomenon of moral luck, but enunciate a theory or conception of praise- and blameworthiness which, if correct, takes the sting from moral luck, i.e., makes it seem unproblematic (despite our ordinary intuitions) that differences in blameworthiness should depend on accidental factors. Let me therefore say a bit more about the ways in which Kantians and utilitarians respectively elude, or seek to elude, the intuitive tangle known as the problem of moral luck.

Kant, as I said, denies the possibility of moral luck, but he has recourse to the metaphysics of noumenal will(ing) in his effort to establish conditions of moral evaluation entirely free from contamination by luck or accident. Nowadays most of us, even many Kantians among us, would tend to resist a noumenal metaphysics—even one in which the noumenal is regarded merely as an inevitable *postulate* of the moral point of view—as the price to be paid for luck-free moral judgment. But if the idea of the noumenal is left to the side, it is difficult to see how the this-worldly determinants of moral evaluation—the factors of circumstance, constitution, and causation that give actuality and substance to moral thought—can reasonably be conceived independently of all luck and accident. And in that case the Kantian still owes us an account of how morality *without* luck is possible.

Utilitarianism, on the other hand, makes the entirely different move of *allowing* for moral luck by conceiving moral predicates in such a way as to make the possibility of moral luck seem less problematic. For the utilitarian, an act is right if it has overall better, happier consequences than any of its alternatives, and an act counts as blameworthy if the further act of blaming it is right in the aforementioned utilitarian terms. (Praiseworthiness is treated in parallel fashion.) In that case the blameworthiness of a person's act amount to nothing more than (reduces to) the fact that blaming it would have better consequences that not blaming it, and since the goodness of consequences

is one of the things in this world and this life which are most subject to luck or accident, such a view makes it easy to see how blameworthiness can depend on factors of luck, if only (or once) the view itself is accepted. Of course, the utilitarian conception of blameworthiness is far from our intuitive ideas about blameworthiness (however muddled the latter may be—for the utilitarian conception is at least not *muddled*). But relative to the acceptance of utilitarian views at least, moral luck becomes fairly unproblematic.

Let us then dwell for a moment on the differences between Kantianism and utilitarianism viewed as attempts to evade the incoherence/paradoxes endemic to our ordinary thinking about moral luck. Kant seeks, in effect, to avoid (our) having to assume the existence of moral luck through the idea of a noumenal realm where human reason and activity are purified of all those empirical factors that so clearly connect with matters of luck or accident. Morality is thereby taken out of the empirical world and construed as pure (inner?) rational willing that bears only an indirect or non-constitutive relation to things and events in the external or phenomenal world. Such a move is clearly reminiscent of and is arguably analogous to the way traditional internalistic/Cartesian epistemology treats epistemic rationality and justification as a function solely of the inner states of the would-be knower.

By contrast, the utilitarian solution to the problem of moral luck is to treat the crucial notions of blameworthiness and culpability as understandable in "externalistic" or naturalistic terms in the light of which the idea of moral luck loses its intuitive unacceptability. We ordinarily think of blameworthiness or culpability as in some sense attaching deeply rather than through luck or accident to moral agents, but a utilitarianism that regards blameworthiness and culpability as matters, respectively, of what it has good/optimal consequences to blame and to treat as guilty considers such matters to be largely *extrinsic* and *accidental* to the agent/act thus deemed blameworthy or culpable. And if, e. g., a person's blameworthiness (for doing x) is a matter of its being the case that blaming the person (for doing x) will have better results for overall human (sentient) happiness than not doing so, then blameworthiness treated as an external property of this kind will also be reducible to non-evaluative empirical terms (assuming that happiness is cashed out in terms of empirically measurable preference satisfaction or pleasure, etc.).

In fact, standard utilitarianism reduces all its ethical vocabulary to supposedly empirical notions/terms—e.g., rightness is just a matter of having consequences for human happiness as good as those of any alternative action—and those who have compared utilitarianism with naturalistic accounts of epistemic terms have seen or come close to seeing that both involve a form of naturalistic *reductionism*. Thus a naturalizing approach that construes epistemic justification in terms of the use of some type of reliable belief-forming process reduces the evaluative epistemic notion of justification to non-

evaluative natural/empirical terms, and this is clearly parallel to the way in which utilitarianism reduces all ethical evaluative notions to natural/empirical ones.[4] But reduction of this sort is not the only procedure open to an epistemological naturalizer, and that fact has, in ways to be discussed shortly, been well known to various recent participants in epistemological debates over the merits of the naturalizing tendency. What has not, however, been known or mentioned is the fact that an ethics aware of the problems of moral luck and conceiving itself as intellectually allied or associated with naturalized epistemology need not assume a utilitarian form. Perhaps it will be easier to see how this is possible, if we now consider an epistemic phenomenon that has received some attention in the literature, but not nearly as much as its ethical analogue: the problem, namely, of epistemic luck.

2. Epistemic Luck and Eliminative Utilitarianism

Many of those who have written on the nature and prospects of naturalized epistemology have pointed out that reductions of epistemic warrant or justification to non-evaluative "natural" concepts introduce a kind of unfairness and/or arbitrariness into attributions of warrant or justification. If, for example, epistemic justification is a matter (very roughly) of arriving at certain beliefs via a reliable (e. g., a typically truth-producing or truth-preserving) cognitive process, then whether or not a given person is (epistemically or rationally) justified in holding a certain belief may depend on factors he has no way of knowing about, on facts beyond his ken and unavailable to him. Thus in two different possible worlds there may be two people making the same inference about their environment on the basis of exactly similar perceptual data; but if the inference is generally reliable in one of the worlds, but not in the other, then on a reductive account of justification, the two individuals differ in the justification of their inferential beliefs. One of them will be justified in believing what he does about his environment, and the other will not be, even though neither has any view about the reliability of the process that underlies his own belief-acquisition. And it seems epistemically unfair and, from an epistemological standpoint, implausible to suppose that the individuals should thus differ in epistemic warrant, when from the standpoint of their mental activity they are (we may assume) exemplifying exactly similar perceptual and inferential processes.

Reductive, naturalizing reliabilist (or related) accounts of epistemic rationality or justification have frequently been objected to in terms like these,[5] and if we may put matters in terms closer to those we have been using here, but quite naturally invoked in regard to the example just mentioned, we can say that the idea that epistemic rationality or justification can depend on factors of luck or accident totally outside the ken of those being epistemologically evaluated is counterintuitive. Just as it is intuitively

repellent to have to suppose that the moral justification of an agent's performing some act can depend on factors of luck or accident outside the ken or control of the agent, it at least initially seems implausible to hold that someone's epistemic justification should depend on factors outside her ken or control. And that explains some of the reluctance to adopt a reliabilist or any similarly externalistic naturalist account of epistemic evaluations.

However, the matter is somewhat more complicated than anti-reliabilists and anti-naturalists seem to have realized. We have seen that our intuitions about moral luck are not all on the side of denying the possibility of such luck, and something similar seems to be the case with regard to epistemic evaluations generally and most particularly with regard to claims about epistemic or rational justification.

Thus imagine a pair of examiners who have just had the unenviable task of telling a dissertation candidate that his dissertation is unacceptable and requires the most extensive revision. We may imagine that at the dissertation interview, the examiners were in substantial agreement about the defects of the dissertation, and that the candidate himself, when told of their objections, admitted their force and validity with a rather dejected air of resignation. At the point, however, when the examiners emerge from this meeting, one of them says to the other: I wonder whether he'll appeal the decision to higher university authorities. And the other immediately counters by pointing out that there is no reason for the candidate to appeal, since he himself agreed with the strong and shared objections of both examiners. Hearing this reply, the examiner making the original conjecture about an appeal may well feel that he was being irrational, possibly somewhat paranoid, to suggest the possibility of an appeal. But what happens if out of the blue, a few weeks later, it turns out that the student has gone on to appeal—totally unreasonably, let us assume—the examiners' decision on his dissertation? I think the examiner who originally conjectured that he might do so will tend to revise his view about the irrationality or paranoid quality of his own original conjecture. He will feel, instead, that the course of events has in some measure (perhaps completely) *vindicated* his original opinion about the real possibility of an appeal.

However, this is a case of rational luck only if we imagine that the conjecturing examiner does not attribute his vindication to some sort of intuitiveness on his part in the original situation where he made his conjecture. If he, or we, imagine that in the original situation he had some clue about the possibility of an appeal from the way the student was behaving (from the peculiar quality of his dejectedness, for example), a clue which he received subliminally but which was nonetheless functioning *as evidence* affecting the conjectures he was willing to make, then we do not have a putative case of rational luck so much as one of subliminal evidence only subsequently recognized as such.

But I don't think there is any need for the conjecturing examiner, or for us, to see his situation, and his subsequent loss of a sense of being slightly paranoid and irrational, as reflecting the existence of subliminal evidence and the subsequent recognition of its presence. The man himself may be led to revise his estimation of his own earlier rationality or irrationality by the mere fact that his conjecture turned out to be correct. His fears may seem to be (somewhat) justified *by subsequent events*, rather than by a subsequently acknowledged earlier bit of evidence, and this description may indeed be the best expression of the man's own sense of how he is (turns out to be) justified (or less unjustified and certainly not paranoid) in his original assumption. Certainly we ourselves, looking at his situation from the outside, have a tendency to withhold the judgment of paranoidness and irrationality with regard to his original conjecture in a way we would not have been inclined to do if we had simply imagined a case in which no appeal was ever lodged. Even assuming that the candidate provided no subliminal clues to his subsequent behavior, our judgment as to (degree of) irrationality and paranoidness is commonly affected by how we imagine things actually turning out, and this provides for a possibility of (epistemic) rational luck that is quite similar to what we said earlier about the common-sense possibilities of moral luck.

On the other hand, something in us recoils at the idea that epistemic or rational justification (or vindication) in cases like that just mentioned can hinge on subsequent events. And so it would seem that as with moral luck, our initial intuitions and judgments in the area of epistemic or rational luck are not entirely consistent with one another; in that case, again as with moral luck, we cannot simply remain on the level of intuitions, but are in need of theory to tell us which intuitions ought to be abandoned. Since *some* intuitions must be dropped, it cannot be a decisive objection to reliabilist and other naturalistic, reductionist accounts of epistemic justification that they go against some of our intuitions. Rather, the question must be whether, with due weight given to overall considerations of intuitiveness, the total theory arrived at by one or another naturalizing tendency or project is superior to our best internalist, anti-naturalistic, anti-reductionist epistemologies, and this question presumably cannot and should not be closed prematurely.

However, epistemological naturalizers need not end up being reductionists. Some, instead, have advocated that epistemology should be replaced by or turn into a form of psychology (or psychology-cum-neuroscience-cum-biology, etc.). And under such a development epistemic evaluations would presumably be eliminated altogether, rather than reduced to other terms. In explaining human thought and action, we might still be interested in questions of reliability and the like, but an eliminative naturalistic epistemology or psychology would have no (further) interest in whether certain cognitive mechanisms yielded or allowed for rational or justified beliefs or inferences.

A thoroughgoing eliminativism would permit of no such epistemic evaluations (let us assume it would eliminate claims about knowledge along with claims about justification) and would presumably hold that there is no good *theoretical* reason to countenance the making of such evaluations (though there might be no reason to forbid them to every-day life).

I have mentioned the possibility of eliminativism, not because I wish to defend it against naturalistic reductionism or anti-reductionism in epistemology, but because its very possibility can give us a clue to some previously unsuspected ways in which ethics can be naturalized. Standard utilitarianism is the only form of naturalized ethics that (to my knowledge) has been recognized as analogous to a naturalizing trend in epistemology. But once we see that eliminativism is a possible direction for the naturalizer in epistemology, we may be encouraged to look for similar possibilities in (naturalizing) ethics.

Utilitarianism as standardly formulated is a reductionistic form of ethical naturalism, but it is also possible for utilitarianism to present itself in an eliminativist mode, and as such it represents a new and, I believe, interesting way in which ethics can seek to naturalize itself. Moreover, once we see how a thoroughly eliminativist ethical naturalism parallel to eliminativist epistemological naturalism is possible, it will be easier to recognize the possibilities of a naturalistically eliminativist virtue ethics, and the possibilities for naturalizing epistemology will in turn be enriched. But let us begin this process with an account of utilitarianism as a form of eliminativist ethical naturalism.

Eliminative utilitarianism stands to familiar reductive utilitarianism as, say, eliminative materialism stands to reductive materialism (I shall not try to respond to those, like Quine, who hold that the theoretical difference between elimination and reduction is nil or negligible). So just as an eliminative materialism claims there are no mental entities or (realized) mental properties, only physical entities and properties functioning in certain complex ways, an eliminative utilitarianism will deny that anything is right or obligatory or intrinsically good, will deny that there are any exemplified ethical (or evaluative) properties or facts. And the reasoning behind such conclusions will in fact, as in other cases where there is a choice to make between reduction and elimination, follow the reasoning for reductive utilitarianism till it reaches an ultimate parting of the ways with it. Like reductionist utilitarians, the eliminative utilitarian will point out the incoherent, or at least intellectually unsatisfactory, thinking that lies, for example, behind our ordinary moral thinking. She will hold that when one properly strips away or otherwise removes the irrational or unjustified elements in our ordinary usage of "right" and "wrong", one will be left with the clear and humanly significant core idea of producing (a net balance of) pleasure or satisfaction for sentient beings. But here, parting ways with ordinary reductive

utilitarianism, the eliminativist will argue that the sheer enormity of the error and confusion in our ordinary beliefs about moral rightness justifies us in claiming that there is no such thing as moral rightness, etc. By contrast, the reductivist utilitarian holds that rightness turns out to be conduciveness to pleasure or desire-satisfaction and that ordinary moral thinking is simply very much in error about (what) rightness (is).

Clearly, the issue here between eliminative and reductive utilitarianism is familiar from our experience of other disputes regarding the respective merits of an eliminative and a reductive approach. And since the most familiar of these, probably, is the long-standing debate between eliminative and reductive materialism, it is perhaps worth mentioning one possible advantage that eliminative utilitarianism (or eliminative ethics more generally) has over eliminative materialism. It is not obvious that sense experience and thinking are merely posited hypothetical entities, and there may be some force to the claim, therefore, that what eliminative materialism eliminates (among other things) are some of the very data which reasonable abductive thinking must seek to account for. But the idea that we immediately experience certain ethical data or properties is, I believe, an (even) harder thesis to defend than what we have just allowed as a possibility in the area of mind and body, and to the extent, therefore, we have better or more immediate knowledge of our own experience(s) than of ethical facts and properties, utilitarian eliminativism may well be in better shape than (or have at least one advantage over) eliminative materialism.

But even apart from this consideration, it seems difficult to find any reason to prefer reductive to eliminative utilitarianism *as a form of naturalism*. If the traditional, reductive utilitarian protests, for example, that the eliminative approach does away with ethics altogether and therefore with utilitarianism itself as a form of ethics, it can be pointed out that this pair of claims is either simply false or altogether begs the issue against eliminativism. After all, eliminative utilitarianism, like eliminative materialism, makes a distinctive claim of its own. It doesn't remain silent or somehow render it impossible for claims to be made, but rather comes out openly and asserts that nothing is good or virtuous or obligatory—or, if you prefer, that "good", etc., do not denote any properties of things. To be sure, problems can arise here about how the negative claims of eliminative utilitarianism are most properly formulated, but there are similar problems about how to express eliminative materialism, and there is no reason, in either case, to say that the eliminativist eliminates the philosophical field he or she is supposed to be working in. Just as the eliminative materialist holds a metaphysical position and is naturally regarded as a materialist of one particular stripe, the eliminative utilitarian subscribes to a particular ethical view, and one, moreover, that is for similar reasons naturally regarded as a form of utilitarianism.

On the other hand, what one may mean by saying that the eliminativist

does away with ethics is that the eliminativist does away with, in the sense of denying the existence of, ethical properties or facts. But even granting that there is a sense in which this charge is true, it hardly seems to constitute an intelligible reason for preferring reductive utilitarianism over eliminative. In the sense presumably intended, the claim that the latter does away with ethical facts or properties simply points out the essential difference between the two forms of utilitarianism. And one cannot treat that as an argument against eliminativism without essentially begging the question against it.

Certainly, if one denies the existence of good things or right acts in favor of the (mere) existence of pleasure and things conductive to it, that may have an effect on those who hear the denial(s), an effect which the reductivist may wish to call bad and that even the eliminativist can say goes against overall human happiness or pleasure. But because the eliminativist may care as much about the advancement of human happiness as any reductionist, he may have equal reason or motive to keep his theoretical views from becoming known. If, as so often happens, a reductive utilitarian can defend the validity of his views while disavowing their practical usefulness and recommending that they remain esoteric, the eliminativist can do something entirely analogous even without making use of specifically ethical or evaluative notions. And so eliminative utilitarianism cannot readily be undercut as a theoretical option by pointing to the consequences of its being adopted by people generally.

The choice between reductionist and eliminationist utilitarianism is not an easy one. And likewise it is unclear for naturalizing epistemology whether it is better to reduce terms like "justified" and risk the consequent unintuitiveness of certain judgments that then have to be made about particular examples[6] or to eliminate such intuitive judgments by the draconian method of forbidding or deeming false all epistemic value claims. At least there is considerable diversity of opinion on this question among naturalizing and even anti-naturalist epistemologists. And by the same token I don't know how to say anything definitive about the choice between reductionistic and eliminativistic utilitarianism. To be sure, accepting the former guarantees a commitment to counter-intuitive moral judgments and, in particular, to the possibility of moral luck, but eliminative utilitarianism makes its own counterintuitive judgments, and in any event both forms of utilitarianism avoid the *inconsistencies* involved in our intuitive judgments concerning moral luck.

However, another way of avoiding the problems of moral luck needs to be mentioned at this point. This alternative ethical approach is very different from any form of utilitarianism, though it resembles eliminative utilitarianism at least to the extent of advocating the elimination of certain ethical properties/epithets from our ongoing ethical theorizing. What I have in mind is a form of virtue ethics that seeks to avoid the paradoxes of moral luck by being *selectively* eliminative in regard to ethical properties/terms. It will turn out that selective eliminativism is also an attractive possibility in

naturalizing epistemology, and our brief sketch of the virtue-ethical approach I want to bring to your attention may help to bring to light some previously neglected possibilities in naturalizing epistemology.

3. Naturalizing Virtue Ethics

Some paradigmatic forms of virtue ethics seem totally opposed to any naturalizing approach to ethical theory. It is sometimes said that an emphasis on inner motivations/dispositions and a primary reliance on aretaic, as opposed to deontic, ethical terms typify any virtue ethics worthy of the name, and given these criteria an ethical view like James Martineau's, which defines a hierarchy of morally better and worse motives and claims that actions are to be evaluated solely in terms of their (previously or independently evaluated) motives, is paradigmatically a form of virtue ethics. Yet its primary and exclusive emphasis on inner motivation seems to ally it with epistemological Cartesianism and the latter's "subjective turn" rather than with the kind of naturalizing epistemology that refuses to base its evaluations solely on accessible and inner mental factors. Indeed, an ethical theory like Martineau's—and I mention this unfamiliar view because it is so simply formulated and such a clear-cut instance of virtue ethics—seems to lie at the opposite end of the ethical spectrum from utilitarianism, and if all virtue ethics had to resemble Martineau's view in the relevant respects, it would make no sense to mine the field of virtue ethics for a naturalizing example.[7]

But whereas Martineau's virtue ethics seems capable of making use, in its distinctive way, of all the main concepts/terms of standard ethical theory, there is another, quite different form of virtue ethics—historically familiar from Spinoza's *Ethics*—which *restricts* its ethical terminology in the light of problems raised by a more extensive ethical vocabulary. Spinoza denies the possibility of metaphysical human freedom and on that basis refuses to allow attributions of moral praise- or blameworthiness into his theoretical account of ethical phenomena. But he is willing to speak of certain character traits as virtues or vices, and as admirable or not admirable, because he assumes we can make sense of these notions independently of any assumptions about metaphysical freedom of will.

For Spinoza some people can be better or more excellent than others in various respects—e.g., one person might be a lovely person, another a vicious human being—though those judged worse in these ways are not thereby be deemed blameworthy or more blameworthy than those judged to be better. The absence of freedom undercuts moral evaluations that inherently assume some sort of metaphysical freedom on the part of human beings, but other sorts of evaluation do not entail such freedom and thus, according to Spinoza, apply to the sort of metaphysically determined but rational creatures we

humans are or can be. A person who frequently turns on people un-expectedly—someone who acts angrily and aggressively toward people, without having been given any provocation—can be regarded as vicious and be avoided as such independently of any commitment to blame the person for being vicious and acting or interacting badly with others (after all, a dog can be called vicious for similar reasons). So Spinoza holds, and we can follow him in holding, that ethical evaluations need not commit us to freedom of will or (therefore) to ascriptions of moral blameworthiness, moral praise-worthiness, or moral responsibility generally.

Nowadays, we are less confident than Spinoza was that causal/metaphysical determinism makes human free will impossible, but we have another motive for wanting to avoid moral/ethical language that commits us to ascriptions of moral praise—and blameworthiness that Spinoza lacked. For we have seen that it is precisely with respect to ascription of blameworthiness and the like that ordinary intuitive thinking ties itself up into knots; the paradoxes of moral luck most closely concern such ascriptions, and so one way to avoid the paradoxes is simply to avoid ascribing blameworthiness, etc., altogether. An ethics of virtue that speaks of admirable and deplorable traits of character and of virtues and vices (or anti-virtues) in the manner indicated by Spinoza can avoid the paradoxes of moral luck by simply eliminating those ethical/moral terms whose ordinary use gives rise to the paradoxes. And this way of dealing with moral luck is quite different both from eliminative and from reductive utilitarianism.

Unlike eliminative utilitarianism, Spinoza-like virtue ethics is only selectively eliminative of moral/ethical concepts/terms, and the concepts/terms it eliminates are (among) those utilitarianism retains, but (re)interprets, reductionistically, in empirical, naturalistic terms. We have thus uncovered three different ways in which naturalizing ethical views can seek to take the sting out of the problem of moral luck. But since it may not yet be clear why I want to hold that virtue ethics of the sort just described should be viewed as a *naturalizing* alternative to utilitarianism, we have some more explaining to do, and the further explanation will help us, in turn, to identify a promising form of naturalizing epistemology that has not yet received the attention it deserves.

One of the most important aspects of naturalizing epistemology has been its typical commitment to externalism in regard to epistemic/evaluative attributions. For the Cartesian epistemologist epistemic rationality and/or justification is a matter of the thoughts, perceptual experiences, and inferences of the would-be knower, and thus concern only the internal mental states of that knower. But an externalist will treat rationality and/or justification as at least partly involving matters external to the mind or subjectivity of the person whose rationality/justification is in question. And a form of externalism like reliabilism with respect to epistemic justification, by making

such justification depend in part on how reliable certain inferential processes actually are in representing our environment to us, makes epistemic justification depend on relations between the mind and the (rest of) the natural world. By contrast, internalism may or may not locate the mind at a point in the natural world but it leaves epistemic justification having nothing to do with (the rest of) the natural world, and this illustrates, I think, the clear sense in which externalism is a typical and exemplary feature of naturalizing epistemology.

But one of the thoughts that help to give rise to the paradoxes of moral luck is our ordinary belief that moral blameworthiness, badness, and goodness are a matter of inner willing or intention, not of possibly accidental and/or unforeseeable extra-subjective effects or circumstances. And to the extent, for example, that a Kantian or intuitionist places a primary emphasis on moral evaluation and sees such evaluation as based in the inner or mental life of rational agents, such an approach to ethics seems highly analogous to Cartesian epistemological internalism, and it is not surprising, therefore, that Kantian epistemology is a paradigmatic (though of course highly distinctive) example of Cartesian epistemological subjectivism.

So an ethical emphasis on the moral as inner parallels Cartesian epistemological internalism, but both these positions emphasize only one side of our ordinary ethical and epistemic evaluations. As we have seen, another part of our thinking seems to want to allow for luck, and both reductionist utilitarianism and reductionist naturalizing epistemology rely on and develop this other side of our ethical and epistemic thinking. Each of the latter focuses on extra-subjective factors that may be thought relevant to ethical or epistemic evaluation(s), and so reductionist naturalizing epistemology and standard reductionist utilitarianism represent antitheses to anti-naturalistic epistemological and ethical internalism.

But there is another way of dealing with the tensions and contradictions in our evaluative thinking than by emphasizing one element of the contradiction or paradox at the expense of the other(s). One may find reasons in the contradiction and paradox to drop the terms that give rise to them, and this, we have seen, is how a naturalistic but eliminative utilitarianism and an eliminative naturalizing epistemology-as-psychology both think we should proceed. Or, as a virtue ethics influenced by Spinoza would urge, we can be selectively eliminative: eliminating the (moral) terms whose dependence on luck seems (to some of our strong intuitions) so objectionable in favor of ethical terms/concepts whose dependence on luck or accident or the unforseen seems much less problematic and which (therefore) are less problematically understood in externalist fashion than such specifically moral terms as "blameworthy," "reprehensible," and "morally good."

Thus a person who was maltreated as a child and who (let us assume) became subsequently incapable of getting along well with others can be

regarded as being the way he is as a result, in large part, of external, unlucky factors outside his control or ken. But that is no reason to deny that the person gets on badly or poorly with others if the latter notion is understood—as it seems quite natural to do—as entirely independent of the supposition that the person in question is morally blameworthy or reprehensible or responsible for being as he is. And so a Spinoza-inspired ethics of virtue can make use of evaluations that clearly and unproblematically depend on factors external to the agent (her will and knowledge).

Consider a further example. One component of being a good father and, as we may now assume, a good mother as well is to be a good provider for one's family. Imagine, therefore, a father who is laid off work after many years on the job, but who, after several weeks of vain searching for another, comparable position, finds such a position through the sheerest luck (imagine, if you will, that thousands of people with his skills have been laid off and that they are in competition for the small number of positions requiring those skills that are available or opening up). In that case, if he takes the position and is again able to provide for his family in the manner to which they had been accustomed, he will eventually be considered to have been a good provider. But if, on the other hand, he had failed to find any good job and were never again able to provide his family with any sort of comfort, then the overall claim that he is or has been a good provider would be irremediably undercut. Circumstances of the sort just described are familiar from the historical example of the Great Depression of the 1930's, and in regard to such a period it seems in no way odd or problematic to claim that factors of luck played a (considerable) role in determining who was and who was not (able to be) a good provider for his family. (Similar points can be made about the notion of taking good care of one's children/spouse/family.)

The people who were unable to find jobs in the Depression through no fault of their own cannot be considered to have been good providers, but in evaluating them thus we need not condemn or blame them, in fact it is difficult to see how anyone *could* blame them for failing, in such circumstances, to provide well for their families. (It is tragic that a whole generation of men and women should so largely have lacked access to this kind of normal human role-ethical attainment, but that is another matter.) And it is the absence of any implication of blameworthiness or reprehensibility here that allows judgments of good providerhood to depend so plausibly and unproblematically on extra-subjective factors external to the (consciousness of) the evaluated agent. So a Spinoza-inspired ethics of virtue (and in the next section we shall see, as is perhaps already obvious, that such an ethics need not be fundamentally egoistic like Spinoza's) removes those ethical concepts that push one toward internalism and relies on other concepts (at least some of) which are understandable in an externalist manner similar to (though not exactly the same as) what we find in reductionistic utilitarianism. And though

such virtue ethics doesn't propose to reduce the terms it uses to purely naturalistic terminology, its freedom from internalism and frequent emphasis on external factors of luck give it a definite and substantial resemblance to naturalizing epistemology.[8] What will make the attribution of naturalizing to virtue ethics more plausible, however, will be the fact that naturalizing epistemology allows for an almost exact analogue of the virtue ethics I have just been discussing.

The dispute in the recent epistemological literature between internalists/anti-naturalists and externalists/naturalists largely centers around terms like "justified" and "warranted" which, as we have seen, have aspects favorable to both sides of the dispute. This can lead one to reject all evaluative epistemic terms and move into pure psychology-biology, etc. But, as we saw with ethics, a selective elimination may be possible in which one rejects the terms which underlie and are the source of recent disputes in favor of other evaluative terms that give rise to no similar difficulties, but are less frequently used by epistemologists than the controversial terms "justified," "vindicated," "warranted," etc.

Instead of using the latter, why shouldn't we talk of cognitive mechanisms or strategies or habits that allow or cause *better or worse adaptation* to a creature's (or a person's or a species's) environment or talk of *better or worse cognitive functioning* (of functioning well or poorly)? Such notions/terms take in extra-subjective factors and are clearly relational and externalist; but unlike the terms that are presently the cause of so much epistemological controversy, these notions wear their relationality on their sleeve. Talk of functioning and of adaptation clearly has to do with more than what is inside the mind (or even the body, in most cases) of the individual whose functioning or adaptation is being evaluated as better or worse. And if internalism has no plausibility whatever in regard to such notions, then there is no reason to deny that luck or accident or the unforeseeable can play a role in their applicability. So if we are selectively eliminative in epistemology, we can undercut the problems of rational or epistemic luck analogously to the way in which virtue ethics can undercut the problems of moral luck.

Thus a naturalizing epistemology that seeks to get beyond the opposing positions in recent disputes about epistemological naturalizing need not go all the way to psychology and biology. Or, since much psychology would be willing to talk about cognitive dysfunction and much biology about maladaptiveness, perhaps I should say, rather, that leaving internalism or naturalizing reductionism for science does not, as one might suppose, force one to give up all epistemic or cognitive evaluation(s). Talk about how well and adaptively one's cognitive mechanisms are functioning seems non-committal on the questions of justification and rationality that have been the focus of recent and traditional epistemological debate. And an epistemology that makes use of the former notions and avoids the latter is indisputably

a version or example of naturalizing epistemology.[9] Since, by dint of its Spinozistic talk of better and worse or more and less excellent, the virtue ethics we have been talking about seems a very close analogue of the just-mentioned form of naturalizing epistemology, it may also now be clearer why I have spoken of such an approach to virtue ethics as naturalized or naturalizing. However, I would like at this point, and finally, to consider one further familiar aspect of the naturalizing tendency in epistemology. It too has an analogue in ethics, and in this case I think ethics, and in particular the virtue ethics we have been describing so sketchily, has a great deal to learn from its counterpart in epistemology.[10]

4. Ethics in Mid-Voyage

In our discussion so far of the naturalizing tendency in epistemology and its ethical counterparts, we have not mentioned what is perhaps the historically most important and currently most familiar aspect of naturalizing epistemology. The idea that in philosophy as in science, we must proceed with our actual beliefs and doubts, rather than base our thinking on some isolated set of pure certainties arrived at through some methodological ideal of absolute indubitability, goes back at least to Peirce and is perhaps most forcefully and famously expressed in Neurath's image of sailors who must repair their ship in mid-voyage.[11] In "Epistemology Naturalized," Quine refers to the Neurath simile in describing what he clearly takes to be an essential element of epistemological naturalizing: the commitment to working what one has and already accepts in order to improve the latter and learn more about the world.[12] But although this methodological conservatism has been typical of inquiry in the natural sciences, it has been notably absent from Cartesian-influenced attempts to reconstruct all knowledge on the basis of idealized and absolute certainties. And according to Quine epistemology should change its stripes and, by adopting a similar attitude of methodological conservatism, become properly scientific and naturalistic.

Note, however, that it is hardly clear that methodologically conservative natural science must ipso facto commit itself to reducing or eliminating certain (ordinary) terms. In letting our critical and inventive faculties loose on the corpus of what we already accept, we needn't perhaps prejudge the issue of whether the progress to be made thereby must inevitably occur through the reduction/elimination of various terms/concepts. So Quine and Peirce and Neurath's sober ideal of epistemology done in mid-voyage seems both more fundamental than and at the same time logically independent of the aspects of naturalizing discussed in earlier sections of the present essay. As such, moreover, it has an analogue in ethics that we have so far left unmentioned. If, given the history of ethical theory, anything has a right to

be considered the most fundamental task of ethics, surely it is the task of showing that justice pays—or, more generally, that morality and a concern to do well by other people can be justified. Typically, attempts to show the validity of morality in this fashion have centered around arguments intended to refute egoism, the attitude of pure or fundamental selfishness. And what seems interesting here in connection with the issue of naturalizing is the resemblance this traditional ethical enterprise bears to Cartesian epistemology.

The Cartesian/Kantian epistemologist starts with the subjective data or consciousness of a single individual in isolation and seeks by argument to bridge the putative gap between the subject and the world. Such epistemology starts with what it takes to be the indubitable or more certain and attempts to justify belief in what it takes to be initially less certain and more in need of epistemic/epistemological validation. As such it clearly falls under Quine's strictures, and even without committing ourselves here either to agreeing or to disagreeing with those strictures, we should at this point be able to recognize the possibility of a similar move in a methodologically naturalizing ethics.

Why should an ethics that is properly sure itself, that wishes to proceed as far as possible in the manner of natural or other science, treat the standpoint of self-concern or self-interest not only as more secure than concern for others, but as a theoretically necessary point of departure for any attempt to justify or vindicate the latter? Why does (basic) concern for others need this kind of justification? Why can't we say, rather, that (almost) everyone in recent ethics has rejected (basic) egoism, and that current ethical theory has a perfect right to assume the validity and ethical justification of concern for other people until and unless our accumulating ethical ideas and theories give us reason to reject this fundamental assumption? There were notable egoistic theories (e.g., Stoicism and Epicureanism) in the ancient world, but that need not and should not disturb our own long-standing and on-going assumption of the validity of (non-egoistic) concern for others, any more than the existence of epistemological skeptics in the ancient world and in some cases more recently is a sufficient reason to go in for Cartesian epistemology. If a proper naturalizing methodology undercuts any need in epistemology to vindicate our beliefs about the world and others on the basis of assumptions about the subject in isolation, then a similarly naturalistic attitude in and toward ethical thought and theory should allow us to accept the ethical validity of a fundamental concern for others without (successfully) attempting to show how concern for others can be based in sheer self-concern or self-interest.[13] Both epistemological Cartesianism and the traditional task of defending altruism to the egoist treat the individual or individualistic standpoint as more fundamental than and necessarily involved in the justification of some larger or more inclusive picture of things, and, again, even if we remain neutral or uncertain about the relevance of scientific methodological conservatism

to these areas of philosophy, we can surely see that ethics contains a plausible analogue to the naturalizing epistemological attitude of a Peirce, or Quine, or Neurath. An ethics or an ethical theory that adopts for itself the same methodological attitude as these philosophers recommend for epistemology—though clearly they wouldn't have wanted to limit its application just to one area of philosophy—deserves to be called a naturalizing or naturalized ethics, and I think ethical theory therefore needs at this point to recognize the possibility of such an approach and consider its merits and demerits. And it may be worth our while to consider in particular whether some philosophers who accept the validity of basic concern for others have recently been too concerned about finding a way to refute egoism.

However, I also want to point out that a methodologically naturalizing approach in ethics is perfectly compatible (at least on the face of it) with the alterations in our ordinary ethical thinking that utilitarianism and Spinozistic virtue ethics urge on us for their respective theoretical reasons. As I mentioned earlier, utilitarians believe various elements in our common-sense moral thinking to be incoherent or without adequate foundation, and it is largely on that basis that they propose some version of the principle of utility as the ultimate criterion of right and virtuous action. But the principle of utility retains the concern for others that is fundamental to common-sense morality; it merely gerrymanders it—in some ways shrinking, in others ways enlarging it—in the light of the difficulties and confusions it claims to find in our ordinary morality. This procedure seems in no way out of keeping with the ethical instantiation of methodological conservatism, and given the obvious connection between utilitarianism and the naturalizing tendency, this should come as no surprise. So utilitarianism in appropriate methodological fashion retains concern for others while rejecting other elements in ordinary moral thought, and a naturalizing methodological approach to ethics generally would then insist that utilitarianism has every reason to retain its commitment to a concern for other people and needn't feel any need to present a refutation of egoism.

By the same token, anyone nowadays who starts with the assumption—either theoretically or in daily living—that we should be concerned with other people and who decides, on the basis of their inherent difficulties, that we should abandon Kantian, common-sense, or utilitarian moral theory for an ethics of virtue, should favor an altruistic over an egoistically Spinozan ethics of virtue. We saw earlier that what distinguishes an ethics of virtue that seeks to handle the paradoxes of moral luck (and other theoretical problems that we have no space to mention here)[14] is the absence of those concepts of morality that give rise to the paradoxes. Such an ethics eliminates the notions of moral blameworthiness or moral goodness in favor of talk about what is better or worse, excellent or poor, admirable or criticizable in human conduct or motivation, and though it is true that Spinoza eliminates the former notions

for the latter in a fundamentally egoistic manner, there is absolutely no reason, given methodological conservatism, why we should follow him in this respect.

An adherent of the naturalizing approach to ethics can therefore claim that we can and should formulate any virtue ethics we may wish to defend without assuming that we need to show how such an ethics can refute egoism. In other words, a virtue ethics that avoids (certain) specifically moral categories can fundamentally deplore attempts to harm or failures of concern for other people, and in the light of methodological conservatism has every right and reason to do so until and unless something turns up to cast substantive doubt on such non-egoistic terms of criticism.

Does that mean that we have no motive to show how other-regarding reasons/virtues can be derived from self-regarding (or non-other-regarding) reasons/virtues unless we reject naturalizing in ethics? Not at all.

To be sure, I have only analogized between naturalizing ethical and epistemological ideas and structures, rather than seeking to defend naturalizing across the board in any definitive way, and anyone who feels a pressing philosophical need to defend ethical concern for others might stand the present discussion on its head and reject ethical naturalism or naturalizing for its inability to comprehend the philosophical force of egoism. But, perhaps more significantly, even a naturalizer in ethics or elsewhere has reason to encourage attempts to argue from non-other-regarding virtues/reasons to other-regarding virtues/reasons or at least to welcome such an argument if one is ever successfully produced. The reason, however, has more to do with the metaphysical/reductionist naturalism mentioned earlier than with doubts about the appropriateness of methodological conservatism in science, epistemology, or even ethics.

If (all) other-regarding virtues/reasons can be derived from non-other-regarding virtues/reasons, then the latter can be reduced to the former and this opens up the possibility of an ethics of reasons or virtues that is fundamentally simpler and more systematically unified that an ethics that cannot effect such reductions. So a reduction of other-regarding values to egoistic or neutral coinage helps to achieve the scientific methodological desideratum of systematic unity in theory. There is thus reason even in a naturalistic or naturalizing ethics to applaud the reduction of the other-regarding to self-regarding terms, if such a thing ever turns out to be possible, but such an attitude is in no way incompatible with the naturalizing methodological belief that we can be and presently are justified in accepting certain other-regarding virtues and/or reasons without being able to produce such a reduction. Compare for example, Smart's version of central-state materialism. On his view we can reduce the mental to the physical or neurophysiological without having to presuppose, as clearly Smart does not, that our beliefs about mental states and processes are made more secure or stand in need of justification by such a unifying philosophical/scientific

reduction. A similar view can be taken of the relation between ethical concern for others and egoism or self-interest.

If what we have been saying above is correct, there have been naturalizing tendencies in ethics—e.g., utilitarianism—for a long time now.[15] But they have not been conscious of themselves in relation to the wider spectrum of naturalizing possibilities explored in this paper. It will be interesting to see how ethics can accommodate itself to the broader picture.

Notes

1. See, e.g., Laurence BonJour, "Externalist Theories of Empirical Knowledge," in *Midwest Studies in Philosophy* V (*Studies in Epistemology*), 1980, pp 53-73; Roderick Firth, "Epistemic Merit, Intrinsic and Instrumental," *Proceedings and Addresses of the American Philosophical Assn* 55, pp. 5-23, and Hilary Kornblith's Introduction to *Naturalizing Epistemology*, Cambridge, Mass.· Bradford/MIT, 1985, pp 1-29

2. Arguments to this effect can be found in Thomas Nagel's "Moral Luck," in *Mortal Questions*, Cambridge University Press, 1979, pp. 24-38, in Samuel Scheffler's *The Rejection of Consequentialism*, Oxford, 1982, and in my "Utilitarian Virtue," in *Midwest Studies in Philosophy* XIII (Ethical Theory), 1989, pp 393 ff. These discussions focus on different aspects of common-sense morality's incoherence or self-conflict.

3. See Nagel, *op. cit.* The first discussion of the problem of moral luck I am aware of occurs in Adam Smith's *Theory of the Moral Sentiments*

4. I am assuming—with Putnam, Hempel, and many others—that the reduction of terms/properties can occur non-analytically, as, for example, when we say that water is nothing more than H2O One can most plausibly equate epistemic justification with the use of some appropriately reliable belief-forming process if one doesn't insist that this equation/reduction obtains *ex vi terminorum* And, similarly, utilitarian reductionism needn't commit any sort of naturalistic/definist fallacy. For further comparison of the role of reductions in science and ethics, see my *From Morality to Virtue*, Oxford University Press forthcoming, Chs. 4, 11, 13.

5. See, e g , Hilary Kornblith, "The Psychological Turn," *Australasian Journal of Philosophy* 60, 1982, esp p 244, and Alvin Goldman, "Strong and Weak Justification," in J. Tomberlin, ed., *Philosophical Perspectives* 2, 1988, esp pp. 64f

6. Note that a reliabilist view of justification will hold that in possible worlds (assuming they are possible) where a solitary person is given some totally illusory set of experiences that is exactly like someone's experiences in our actual world the former will lack the latter's *justification(s)* for belief in external things, not merely her *knowledge* of such things.

7. For Martineau's views, see his *Types of Ethical Theory*, 2 volumes, 3rd edit., 1891 Given Martineau's comparative obscurity, it is perhaps worth noting that Sidgwick devoted more space in *Methods of Ethics* to Martineau's ideas than to those of any other of his contemporaries.

8. Cf Ernest Sosa, "The Raft and the Pyramid. Coherence versus Foundations in the Theory of Knowledge," in *Midwest Studies in Philosophy* V, esp. p. 23.

9. Unless it brings in God's purposes as the basis for understanding "functioning well," etc. In "Positive Epistemic Status," *Philosophical Perspectives* 2, 1988, pp. 1-50, Alvin Plantinga makes this sort of notion central to his epistemological view

without committing himself on the necessity of relying on theism. Plantinga also doesn't point out the analogy with Spinoza's ethics, and I should perhaps mention that non-egoistic naturalistic virtue ethics need have nothing to do with theism or Spinoza's pantheism

10. The just-sketched form of virtue ethics is described and defended at length in *From Morality to Virtue*

11 See Otto Neurath, "Protokollsaetze," *Erkenntnis* 3, 1932, p. 206; and C. S Peirce, "The Fixation of Belief," in J Buchler, ed., *The Philosophical Writings of Peirce*, N. Y.: Dover, 1955.

12 See Kornblith, ed., *Naturalizing Epistemology*, pp 15-29.

13. This point is, in effect, a naturalized version of Prichard's intuitionistic thesis that other-regarding moral claims present themselves as binding on us quite independently of any eudaimonistic assumptions we may or may not make. See his "Does Moral Philosophy Rest on a Mistake?" in *Moral Obligation*, Oxford, 1949.

14. See the works referred to in footnote 2

15. Some recent naturalizing trends in ethics that I have not mentioned here include: the use of ethical claims as (best) explanations of historical phenomena by Nicholas Sturgeon and others; attempts by Georges Rey, Daniel Dennett and others to reduce moral-psychological notions like weakness of will and free will to computational/naturalistic categories; and, of course, continuing attempts to give naturalistic *analyses* of ethical terms

I would like to thank Michael Devitt for extremely helpful comments on the present paper.